



Service Oriented Architecture

Version 2.3

Updated
24 May 2025

Lecture Notes, Discussion Points and References

Table of Contents

SOA Syllabus	6
Unit 1: Introduction to Service-Oriented Architecture	6
Unit 2: SOA Design and Modeling	6
Unit 3: SOA Implementation Technologies	6
Unit 4: Security and Governance in SOA	7
Unit 5: SOA Emerging Trends	7
SOA Lab Exercises	8
Exercise 1: Overview of Service-Oriented Architecture	8
Steps / Tasks	8
Exercise 2: Principles and Concepts of SOA	9
Steps / Tasks	9
Exercise 3: Contemporary Trends in SOA	11
Steps / Tasks	11
Exercise 4: Artificial Intelligence (AI) and Machine Learning (ML) in SOA:	12
Steps / Tasks	12
Reference Books	14
Terminology	15
Unit-1 Introduction to Service Oriented Architecture	16
Motivations of SOA	16
What is SOA	19
Key Characteristics of SOA	20
SOA Components	21
Evolution and Historical Context of SOA	22
Early Web Services	23
E-commerce Integration - Amazon.com	24
Travel Booking System - Expedia.com	25
Enterprise Resource Planning - SAP ERP	26
Emergence of SOA Standards	26
AI-Driven Services: A Modern Extension of SOA	28
Benefits and Challenges of SOA	30
Benefits of SOA	30
Contemporary Trends of SOA	32
Cloud Computing and SOA	33
Serverless Computing and SOA	33
Unit-2 SOA Design Principles and Modeling	35
Service Design Principles and Patterns	36
1. Service Coupling	36
Definition: Service Coupling refers to the degree of interdependence between any two	

business processes or services within a system.	36
Preferable State: In SOA, weak coupling is preferred, indicating lower dependency for increased flexibility, scalability and maintainability.	36
2. Service Cohesion	36
3. Client - Server Architecture	37
4. IP Address and DNS	38
5. Forward Proxy and Reverse Proxy	39
6. GraphQL vs REST API	40
7. SQL vs No SQL	41
8. Vertical Scaling vs Horizontal Scaling	43
Vertical Scaling (Scaling Up)	44
Horizontal Scaling (Scaling Out)	44
9. Sharding	45
Sharding ≈ Horizontal Partitioning	46
10. Caching	46
Why Caching Matters in SOA	47
11. Idempotency	48
How It Works in SOA Services	49
12. Circuit Breakers	50
Unit-3 SOA Implementation Techniques	51
Web Services Standards	51
Simple Object Access Protocol	51
Representational State Transfer (REST)	51
Graph QL	52
Examples and Code Snippets	52
Containerization and Orchestration	53
Containerization	53
Docker Container	54
Kubernetes Orchestration	55
Event-Driven Architecture	56
Event Sourcing	56
Event-Driven Messaging Systems	57
API Management and Governance	57
API Design Principles	58
Rate Limiting and Quotas	58
Unit-4 Security and Governance in SOA	60
Security Considerations in SOA	60
Understanding Threat Models	60
Key Security Risks in SOA (Microservices) and its mitigation	61
Security Risks in SOA for Healthcare Apps	63
Security Design Patterns:	66
SOA API Security	66
Unit-5 SOA Emerging Trends	68

Serverless Computing	68
Function-as-a-Service	68
Key Features	69
Operational Characteristics of Serverless Computing	69
Challenges	69
Introduction to AI and ML in SOA	70
Intelligent Agents	71
Characteristics	71
Example Use Case	72
Natural Language Processing (NLP)	72
Key Applications	72
Example Use Case	72
Lab Exercises - Solution	74
Exercise 1: Overview of SOA: Implement a REST Web Service	74
REST Web Service - Python Implementation (GET and POST Methods)	74
REST Web Service - Spring Boot (Java) Implementation	76
Exercise 2: Principles and Concepts of SOA	79
Pub-Sub: Demonstrate a Publisher-Subscriber message exchange using RabbitMQ.	79
RabbitMQ tutorial - "Hello world!"	80
Exercise 3: Demonstrate a Content Delivery Network (CDN)	80
Design a simple Content Delivery Network (CDN) using Python with focus on distributing content efficiently to users from multiple edge servers	81
Exercise 4: Build a AI-driven Customer Sentiment analysis service	82
Design a simple AI driven Customer Sentiment analysis service using ML models and integrate it into a SOA application.	83
Exercise 5: Contemporary Trends in SOA	86
Reference Articles	92

SOA Syllabus

Revision 2.0, Mar 2025 (Updates in **Bold**)

Unit 1: Introduction to Service-Oriented Architecture	
1.1	Overview of Service-Oriented Architecture - Idea of a Service, Key Characteristics, Historical Context, AI-Driven Services
1.2	Principles and Concepts of SOA - Service Loose Coupling, Service Reusability, Service Abstraction, AI Considerations
1.3	Evolution and History of SOA - Early Web Services, Emergence of SOA Standards, Transition to Microservices & Containerization, DevOps and MLOps, Role of Cloud Providers & AI Services:
1.4	Benefits and Challenges of SOA - Business Agility, Interoperability, Challenges in Implementation, AI Integration Challenges
1.5	Contemporary Trends in SOA - Microservices Architecture, Cloud Computing and SOA - Serverless Computing and SOA, AI/ML in the Service Ecosystem
Unit 2: SOA Design and Modeling	
2.1	Service Design Principles and Patterns - Service Cohesion, Service Coupling, Client-Server architecture, Service IP and DNS
2.2	Service Connection and Data Management - Proxy, GraphQL, RESTAPI, SQL, NoSQL, Idempotency
2.3	Designing for Scalability and Resilience - Load Balancing, Sharding, Partitioning, Fault Tolerance, Circuit Breaker Pattern,
Unit 3: SOA Implementation Technologies	
3.1	Web Services Standards - Simple Object Access Protocol (SOAP), Representational State Transfer (REST) - GraphQL, gRPC
3.2	Microservices Architecture and its Relationship with SOA - Decentralized Data Management, Independent Deployment, MLOps & Microservices: - Infrastructure Automation, Automated Model Deployment
3.3	Containerization and Orchestration

	<ul style="list-style-type: none"> - Docker Container, Kubernetes Orchestration, Specialized AI/ML Orchestration, Service Mesh Technologies, Observing AI Microservices
3.4	<ul style="list-style-type: none"> Event-Driven Architecture (EDA) and SOA - Event-Driven Messaging Systems, Pub/Sub patterns, Streaming Pipelines for AI
3.5	<ul style="list-style-type: none"> API Management and Governance - API Design Principles, AI-Specific API Considerations - Rate Limiting and Quotas, AI Endpoint Limits
Unit 4: Security and Governance in SOA	
4.1	<ul style="list-style-type: none"> Security Considerations in SOA - Understanding Threat Models, Common Security Risks in SOA Threats in AI-Driven Services - Security Design Patterns, Zero Trust for distributed microservices (including AI endpoints)
4.2	<ul style="list-style-type: none"> API Security Best Practices - Securing RESTful APIs, Web Service Security Standards (WS-Security) - Securing AI Inference APIs
4.3	<ul style="list-style-type: none"> JSON-based Security - SAML Assertions and Assertions Consumers, Modern Alternatives with JWT/OAuth 2.0 vs. SAML usage in microservices and AI service endpoints
Unit 5: SOA Emerging Trends	
5.1	<ul style="list-style-type: none"> Serverless Computing and its Impact on SOA - Function-as-a-Service (FaaS), Event-Driven Architectures, AI Use Cases in Serverless - Operational Characteristics, Observability in AI-Driven Serverless
5.2	<ul style="list-style-type: none"> Artificial Intelligence (AI) and Machine Learning (ML) in SOA - Intelligent Agents, Predictive Analytics, Natural Language Processing (NLP), AI Orchestration and Workflow
5.3	<ul style="list-style-type: none"> Edge Computing and SOA Integration, - Edge Gateway Architectures, Low-Latency Data Processing, Offline Capabilities, AI Workloads at the Edge

SOA Lab Exercises

Exercise 1: Overview of Service-Oriented Architecture

Objective:

Introduce students to SOA basics by creating a simple, containerized web service that clients can consume.

Steps / Tasks

1. Set Up a Basic Service

- Choose a framework (Flask/Python, Spring Boot/Java, Express/Node.js).
- Create a simple endpoint (e.g., `/products` or `/users`) that returns or manipulates data (e.g., CRUD operations in-memory or in a small database like SQLite).

2. Containerization

- Create a `Dockerfile` for your service.
- Use Docker to containerize and run your service locally (e.g., `docker build`, `docker run`).

3. Client Consumption

- Write a simple client script or another microservice to call your web service.
- Demonstrate basic operations (GET, POST, PUT, DELETE).

4. Documentation and Testing

- Produce a small `OpenAPI/Swagger` specification to define your API.
- Use a tool like `Postman` or `cURL` to test endpoints.

Key Learning Points:

- Foundational SOA concepts (service exposure, discoverability, loose coupling).
- Introduction to containerization for easy deployment and scaling.

Exercise 2: Principles and Concepts of SOA

Objective:

Implement a loosely coupled service architecture using **asynchronous messaging**.

Steps / Tasks

1. Choose a Message Broker

- Use **RabbitMQ**, **Apache Kafka**, or **ActiveMQ**.
- Explain how messaging decouples the producer from the consumer.

2. Publisher Service

- Create a simple service that sends messages (e.g., JSON payload) to the broker whenever an event occurs (e.g., new order placed, data update).
- Containerize it if desired (using Docker) for consistency.

3. Consumer Service

- Implement a separate service that subscribes to the broker and processes messages independently (e.g., logs them, stores them in a DB, triggers a workflow).
- Ensure no direct coupling between publisher and consumer beyond the message format.

4. Observability

- Introduce logging and monitoring for your publisher and consumer services (e.g., using **Elastic Stack**, **Prometheus**, or built-in broker metrics).

Key Learning Points:

- Asynchronous communication for **loose coupling**.
- Event-driven design principles and decoupled service interactions.
- Understanding of **microservices** patterns.

Dept. of ISE, BMSCE, 2025

Exercise 3: Contemporary Trends in SOA

Objective:

Explore Serverless Computing and integrate it into an existing SOA/microservice ecosystem.

Steps / Tasks

1. Set Up a Simple Serverless Function
 - Pick an open source serverless platform like (**OpenFaaS, Apache OpenWhisk, Knative, Kubeless, Fission, or KEDA**).
 - Create a function that performs a specific task, e.g., image resizing, simple text processing, or a quick calculation.
2. Expose the Function as a REST Endpoint
 - Use API Gateway (AWS), Azure Function's HTTP trigger, or Cloud Functions' HTTPS endpoint to make your function externally callable.
 - Verify your function can accept inputs and return outputs.
3. Integration with Other Services
 - Invoke your serverless function from a previously created microservice or a simple client.
 - Demonstrate that the function can be part of a broader SOA. For example, upload an image via a REST service, which triggers the serverless function to resize it and then store the result in a storage service.
4. Observability and Cost Monitoring (Optional Enhancement)
 - Show how to monitor invocation counts, latency and cost metrics for your serverless function.
 - Highlight the ephemeral nature of serverless (cold starts, concurrency limits, etc.).

Key Learning Points:

- Basics of Function-as-a-Service (FaaS).
- Serverless integration with existing services for scalability and event-driven operations.

- Challenges like cold starts, limited runtime environment, debugging in a serverless context.

Exercise 4: Artificial Intelligence (AI) and Machine Learning (ML) in SOA:

Objective:

Build and integrate a **simple AI-driven service** (e.g., classification, sentiment analysis, or basic prediction) within an SOA-based architecture.

Steps / Tasks

1. Develop/Obtain a ML Model

- Use a small classification model (e.g., scikit-learn or TensorFlow).
- Pre-train or load a pretrained model (e.g., for text sentiment or an image classification dataset like MNIST).

2. Build a Service for ML Inference

- Wrap the model in a REST endpoint (Flask, FastAPI, or any framework).
- Accept input data (text, image, numeric features) and return inference results.

3. Containerize

- Package the model service with **Docker** for easy deployment.
- Show how the model can be scaled independently, if needed.

4. Integration and Testing

- Integrate your AI service with a front-end client or another microservice.
For instance, the client sends text or image data and the AI service returns a prediction.
- Demonstrate how updates to the model (newer version, better accuracy) can be swapped in with minimal disruption to the rest of the SOA.

5. Expand with MLOps

- Briefly mention or demonstrate how to track model versions, use a model registry (e.g., MLflow).
- Automated testing: ensure new model versions do not break the interface or degrade performance.

Key Learning Points:

- Basic AI model serving in a **service-oriented** environment.
- Handling model versioning, data input/output formats and performance considerations.
- Implementation of MLOps

Dept. of ISE, BMSCE, 2025

Reference Books

1. Service-Oriented Architecture: Concepts, Technology and Design by Thomas Erl
2. Building Microservices" by Sam Newman
3. Microservices Patterns: With examples in Java by Chris Richardson
4. SOA Security by Ramarao Kanneganti and Prasad Chodavarapu
- 5. Designing Data-Intensive Applications – Martin Kleppmann**
- 6. Practical MLOps – Noah Gift, Alfredo Deza**
- 7. Kubernetes Patterns – Bilgin Ibryam, Roland Huß**

Dept. of ISE, BMSCE, 2025

Terminology

Enterprise	Enterprise refers to an organization or a business
Service	A basic granular unit of a system that provide a specific function
Architecture	An organization or design pattern of an software system
IT Systems	Refers to hardware and software components of an Enterprise
Applications	A software designed to specific functions or services
Web Service	An HTTP based application used over internet
Security	Refer to protecting user, data, infrastructure and applications of an enterprise

Unit-1 Introduction to Service Oriented Architecture

1.1	Overview of Service-Oriented Architecture - Why learn SOA, Key Characteristics, Historical Context, AI-Driven Services
-----	--

Motivations of SOA

Consider **modern applications** that we use in our daily lives:

- Streaming - JioHotstar, Netflix, Prime Video
- Payment (NCPI: BHIM, Paytm, Google Pay, etc)
- Cab Booking applications (Ola, Uber, Rapido, Namma Yatri, etc)
- e-Commerce (Amazon, Flipkart)
- Quick Commerce (Zepto, Blinkit)
- Navigation (Google Maps, Open Street Maps)
- Food Delivery (Swiggy, Zomato, etc)
- Social Media (Instagram, Twitter)
- Communication (WhatsApp, Telegram, Jabber)

Notice the number of users watching the live stream on the JioHotstar platform: It's around **66 crore** viewers. In other words, massively scalable system design.

IND vs AUS, Champions Trophy 2025: Jio Hotstar live streaming viewership surges over 66.9 cr as India wins semi-finals

IND vs AUS, Champions Trophy 2025: India won the semi-final against Australia scoring 267 with 4 wickets and 11 balls remaining on Tuesday. Viewership on Jio Hotstar surged to over 66.9 crore as of the end of the cricket match.

Livemint

Updated • 4 Mar 2025, 09:40 PM IST



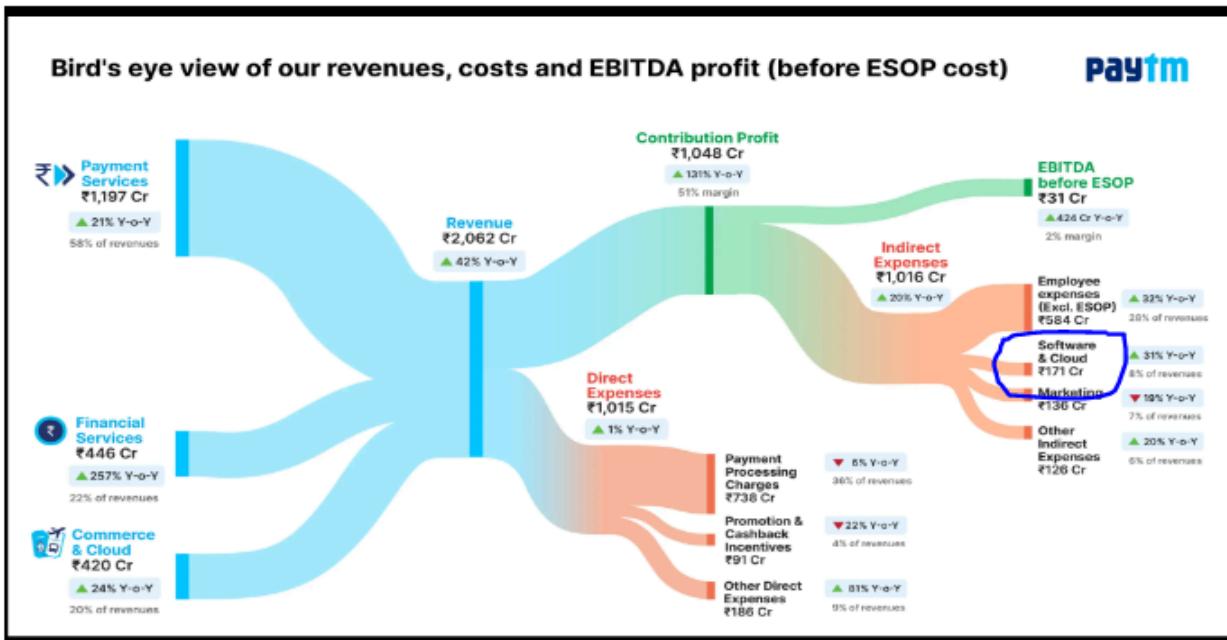
Notice the number of UPI transactions happening on the NPCI platform: around **500 million in a day**. In other words, **highly scalable, resilient and secure** system design.

Day	Volume (in Mn.)
February 1, 2025	604.75
February 2, 2025	591.60
February 3, 2025	582.39
February 4, 2025	576.42
February 5, 2025	585.68
February 6, 2025	579.32

NPCI Transactions: <https://www.npci.org.in/statistics/monthly-metrics>

Challenges of building such applications

- **Scalability** - very high number of users (thousands, millions of active users, transactions)
- **Multiple and Heterogeneous components** - each application has many components interconnected (front-end, backend, databases) and each component is designed differently and coded in different languages (e.g. Java, Go-lang, JS, Python, etc)
- **Security** - authentication (e.g. login), protecting data (encryption), data privacy
- **High Availability and Business Continuity** - making applications available with minimal down time.
- **Cost and Operational Expenses (OpEx)** -
 - Development Cost - developing new products and features
 - Deployment Cost - high cost of running these applications in data centers and cloud.
 - Paytm OpEx: Software, Cloud and Data Center costs were ₹171 Cr, up 31% YoY in Feb 2023, ₹188 Cr in Mar 2023



Source: <https://paytm.com/blog/investor-relations/how-paytm-achieved-operational-profitability/>

What is SOA

Home Work References

1. Read about original publication on SOA:
<https://www.opengroup.org/soa/source-book/soa/index.htm>
2. Read Martin Fowler (ThoughtWorks) take on SOA for an article in 2005
<https://martinfowler.com/bliki/ServiceOrientedAmbiguity.html>

Modern take on SOA (Inspired by Martin Fowler article)

- SOA is about **exposing software and its functions through web services**.
- SOA is also defined as “big apps,” **focusing on small core services** plus aggregator UIs for business and AI tasks.
- SOA is also referred to as a **universal messaging backbone**—HTTP, Kafka, etc, to share data across systems..
- SOA is also referred to as **for asynchronous (non-blocking) data flows**, letting AI workflows happen behind the scenes without blocking.
- Overall, **SOA**— is about **building flexible, maintainable systems, often mixing microservices, serverless and AI**.

Formal Definition

- **SOA refers to** → Service Oriented Architecture
 - Service-Oriented Architecture (SOA) is an architectural style that supports **service-orientation**.
 - **It's a Design Pattern:** a way to build modern complex applications using **granular, reusable services**.
 - **It's an approach** to build software systems that are based on **distributed systems**.
 - It's an approach to build software systems based on **loosely coupled service** components
- A **service**:

- Is a logical representation of a repeatable business activity that has a specified outcome (e.g., check customer credit, provide weather data, consolidate drilling reports)
- Is self-contained
- *May be* composed of other services
- Is a “black box” to consumers of the service
- **Idea of a Service** - A service is defined as a specific **granular, functional, self-contained, reusable** component or code consumed by other services or applications (e.g. Login Service, Order History, Map APIs)
 - **Service interface** - provides interface to invoke a service and define formats to pass and receive data from a service. For example, users of RESTAPI for request and response, XML, JSON for sending and receiving data from service.
 - **Service is technology independent and interoperable** - consumers of the service can invoke the service on any hardware or software platform or code. For example, a Cab booking app running in AWS can invoke Google Map API services to from source to destination
 - **Service is discoverable** - consumers of the service can easily detect the purpose and use of the service. For example, an E-Commerce app can discover various payment methods.
 - **Service is stateless** - a service doesn't maintain any specific state of a service call. For example, a QR scanning service takes a QR code, just returns the value of code and doesn't maintain any other context of service invocation.

Key Characteristics of SOA

Service-Oriented Architecture (SOA) is defined by several key characteristics that shape its design and implementation. These characteristics include:

1. **Loosely Coupled**: SOA promotes loose coupling between software components, allowing them to interact independently without tight dependencies. This enables

flexibility and agility in system design, as services can be modified or replaced without impacting other components.

2. Interoperable: SOA facilitates interoperability between heterogeneous systems and technologies. By adhering to open standards and protocols, such as XML, SOAP and REST, services can communicate seamlessly across different platforms and programming languages.

3. Flexible: SOA is inherently flexible, allowing for the composition and recombination of services to meet changing business requirements. Services can be combined and orchestrated in various ways to create new functionalities, enabling organizations to adapt to evolving needs.

4. Scalable: SOA provides scalability by distributing functionality across multiple services, each capable of running independently and horizontally scaling to accommodate increased demand. This allows systems to handle varying workloads and scale resources efficiently.

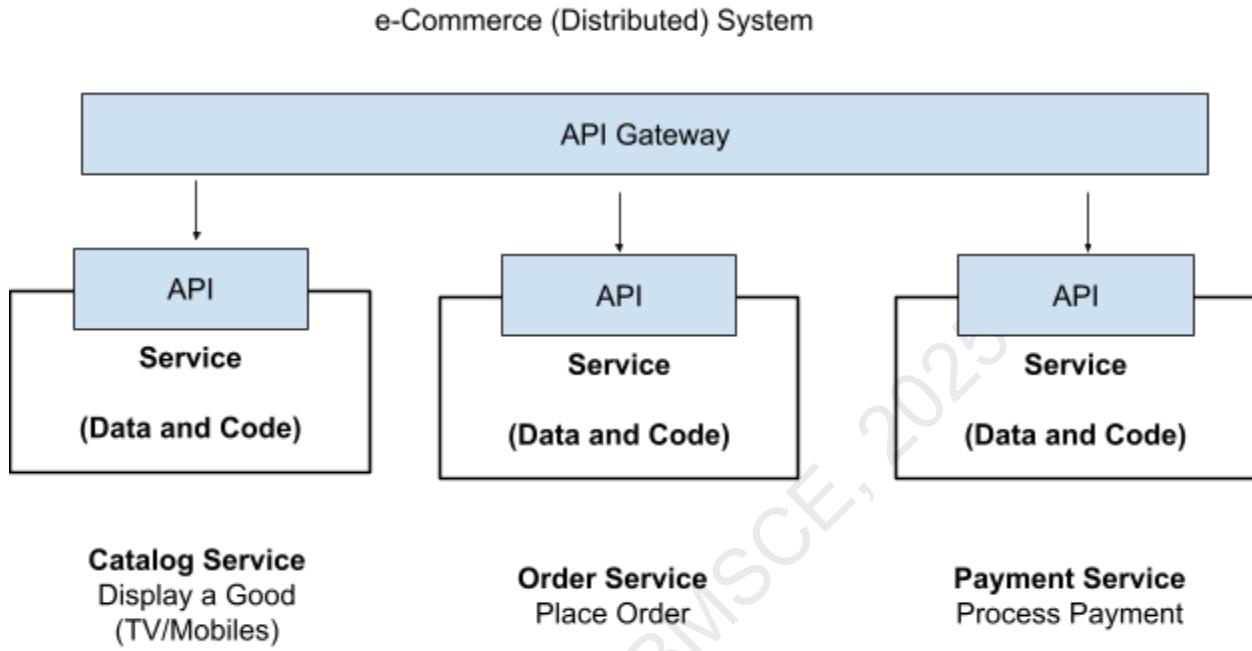
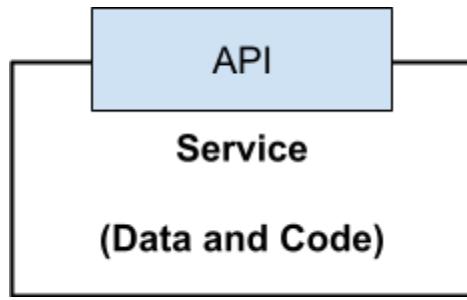
5. Stateless: SOA promotes statelessness, where services do not maintain session state between requests. This enhances scalability and fault tolerance by allowing services to handle each request independently, without relying on previous interactions.

These key characteristics of SOA—loose coupling, interoperability, flexibility, scalability and statelessness—lay the foundation for building resilient, adaptable and efficient **complex and distributed** software systems.

SOA Components

The basic unit of Service Oriented Architecture is a **Service**

- It's a self-contained software system
- Accessible via APIs



Evolution and Historical Context of SOA

Late 1990s - Early 2000s: Emergence of Web Services

- Key Events:
 - Rise of Internet technologies - Dot Com Boom! Web based growth
 - The term first appeared in 1998
 - Need for interoperability between systems
- Technologies:
 - Introduction of SOAP (Simple Object Access Protocol) and WSDL (Web Services Description Language)

Early to Mid-2000s: Paradigm Shift with SOA

- Key Events:
 - Transition towards modular, loosely coupled architectures
 - Emergence of SOA as a new architectural paradigm

- *Technologies:*

- Adoption of SOA principles in enterprise integration

Standardization Efforts and Industry Consortia

- *Key Events:*

- Development of SOA-related standards and specifications
- Role of organizations like W3C and OASIS

- *Technologies:*

- Definition of XML, SOAP, WSDL and WS-* standards

Challenges and Criticisms in SOA Adoption

- *Key Challenges:*

- Complexity in service design and governance
- Cultural resistance to change
- Concerns about ROI

- *Critiques:*

- Effectiveness of SOA in delivering promised benefits

Legacy and Impact

- *Key Influences:*

- Evolution towards microservices and cloud-native architectures
- **Serverless** and **API-first** approaches also align with SOA's idea of exposing discrete functionality.
- Continued relevance in modern software design practices
- The concept of **service boundaries** remains crucial, especially in modern **AI/ML** model-serving scenarios.

- **Continued Relevance:**

- Even as technology stacks shift (REST, GraphQL, gRPC), SOA's core vision of **interoperable services** still underpins **enterprise integration** and **distributed architectures**.

Early Web Services

Here are use cases of SOA in early web services:

- **E-commerce Integration:** Early web services facilitated the integration of e-commerce platforms with payment gateways, enabling secure transactions and real-time order processing.

- **Travel Booking Systems:** Travel agencies use web services to integrate with airline reservation systems, allowing customers to search for flights, book tickets and receive real-time updates on flight availability.
- **Enterprise Resource Planning (ERP):** Organizations utilise web services to integrate ERP systems with third-party vendors for tasks such as inventory management, supply chain optimization and financial reporting.

E-commerce Integration - Amazon.com

W | YOUR ACCOUNT | HELP

WELCOME BOOKS MUSIC VIDEO TOYS & GAMES ELECTRONICS e-CARDS AUCTIONS

HOW TO ORDER | OUR GUARANTEE | SITE GUIDE

Friday, August 27, 1999

SEARCH

All Products

WELCOME TO amazon.com

Hello! Shopping at Amazon.com is 100% secure--guaranteed.
Already a customer? [Sign in](#).

BROWSE

- [Books](#)
[Bestsellers](#), [Computers](#), [Kids](#), [Business](#)...
- [Music](#)
[New Releases](#), [Top Sellers](#), [Classical](#), [Soundtracks](#)...
- [Video](#)
[DVDs](#), [Top Sellers](#), [New Releases](#), [Kids & Family](#)...
- [Auctions](#)
[Books](#), [CDs](#), [Photo Equipment](#), [Collectibles](#)...
- [Electronics](#)
[Cameras](#), [Computer Add-Ons](#), [DVD Players](#)...
- [Toys & Games](#)

In Books
London Calling

Wrapping up an ambitious reissue program, the great Penelope Fitzgerald delivers not one but two classics in miniature. *The Golden Child* is the British author's first novel, in which an ancient mummy ignites a curatorial ruckus, while *A Freddie's* revolves around a drama academy and its deeply amusing student body. Go to [Books](#)

In Electronics
MiniDisc the Magnificent

Sony's [Bundle5 MiniDisc package](#) harmonizes two separate products--a MiniDisc home recording deck and a portable MiniDisc player--so you can record, edit, and title your own discs from both analog and digital sources, then enjoy them on the go. Go to [Electronics](#)

In Toys & Games
Pack with Pride

What are they listening to in [Austin, Texas](#)? Explore Amazon.com [Purchase Circles](#).

Amazon.com 100 Hot Books

Updated Hourly

1. [Harry Potter and the Prisoner of Azkaban](#) by J. K. Rowling
2. [Harry Potter and the Sorcerer's Stone](#) by J. K. Rowling
3. [Harry Potter and the Chamber of Secrets](#) by J. K. Rowling

► [More Hot Books](#)

- **Use Case Scenario:** In the mid-90s, Amazon.com began its journey as an online bookstore, but it quickly evolved into a platform selling a wide range of products. To handle the complex nature of its e-commerce operations, Amazon likely employed service-oriented principles, even before SOA was formally recognized.
- **SOA Characteristics:** Amazon's platform likely utilised loosely coupled services to handle various aspects of its operations, such as inventory management, order processing, payment processing and customer relationship management (CRM).

These services would have communicated with each other using standardized protocols such as HTTP and XML

Travel Booking System - Expedia.com

- **Use Case Scenario:** Expedia, founded in the late 1990s, revolutionized the travel industry by offering an online platform for booking flights, hotels, rental cars and vacation packages. To provide a seamless booking experience to its users, Expedia likely employed SOA principles in its system architecture.
- **SOA Characteristics:** Expedia's platform likely consisted of various services responsible for different functions, such as flight search, hotel booking, payment processing and itinerary management. These services would have been loosely

coupled, allowing for flexibility and scalability. For example, the flight search service could communicate with airline reservation systems via standardized interfaces, such as SOAP or XML over HTTP.

Enterprise Resource Planning - SAP ERP

- **Use Case Scenario:** SAP, a leading provider of ERP software, has been in the industry since the early 1970s. While its early systems may not have fully embraced SOA principles, SAP likely transitioned towards a more service-oriented approach in the late 1990s and early 2000s to address the growing complexity of enterprise operations.
- **SOA Characteristics:** SAP's ERP system would consist of various modules or services responsible for different business functions, such as finance, human resources, supply chain management and customer relationship management. These services would have been designed to be reusable and interoperable, allowing organizations to customize their ERP implementations based on their specific needs.

Emergence of SOA Standards

The emergence of SOA standards occurred as organizations sought more structured and scalable approaches to integrating systems and building software applications.

Key Standards	Original Role	Modern View
SOAP (Simple Object Access Protocol): A protocol for exchanging structured information in the implementation of web services.	Defined message formats, error handling and communication protocols for web services .	Still used in many enterprise environments, especially where WS-* standards (e.g., WS-Security) are critical.

Key Standards	Original Role	Modern View
WSDL (Web Services Description Language): A standard for describing the functionality of web services, facilitating their discovery and invocation.	Provided a contract-first approach to describe web service interfaces and operations, facilitating discovery and invocation in SOAP-based contexts.	Inspired later mechanisms in RESTful environments, such as OpenAPI/Swagger for describing REST APIs.
XML (eXtensible Markup Language): A markup language used for encoding documents in a format that is both human-readable and machine-readable	Became the go-to format for structuring data in SOAP messages and configuration.	Shift to JSON: Over time, many services adopted JSON due to its lighter weight and better fit for web-based communication.
REST (Representational State Transfer) & JSON	Gained popularity as a simpler alternative to SOAP, using HTTP verbs (GET, POST, etc.) and JSON payloads	Encouraged stateless interactions and easier integration, forming the backbone of many microservices and modern APIs.
gRPC & Protocol Buffers	A high-performance, language-agnostic RPC framework that uses	Well-suited for low-latency or AI/ML scenarios, often used for microservices at scale.

Key Standards	Original Role	Modern View
	Protocol Buffers for data serialization.	

AI-Driven Services: A Modern Extension of SOA

Modern **Service-Oriented Architecture (SOA)** isn't limited to traditional web services that simply exchange XML or JSON over HTTP. With the rise of **artificial intelligence (AI) and machine learning (ML)**, organizations increasingly expose **trained models** as **independently deployable services**—often referred to as "**Model as a Service**" (MaaS). This trend significantly extends the original SOA concepts by emphasizing:

1. Data-Centric Interfaces

- Traditional SOA focuses on functional operations (e.g., "create order," "update account").
- **AI-driven** endpoints often revolve around **model inference**: accepting data (e.g., images, text) and returning predictions (e.g., classifications, recommendations).
- Data formats can be more complex—images, audio, time series—requiring careful consideration of **serialization** (JSON, Protocol Buffers, etc.).

2. Loose Coupling for Continuous Model Updates

- A key SOA principle is **loose coupling**—ensuring that changes in one service do not break others.
- When models are treated as **services**, updates (like retraining or swapping a model) can happen **independently** without disrupting downstream clients.
- This aligns with **DevOps/MLOps** practices, where continuous integration and deployment (CI/CD) extend to **model retraining** and **re-deployment**.

3. Scalability and High-Performance Requirements

- AI inference can be **resource-intensive**, especially for large models (e.g., deep learning) that may run on GPUs.
- **SOA** must account for dynamic **scaling** and specialized hardware needs to handle fluctuating inference workloads.
- Auto Scaling microservices (on Kubernetes, for instance) or employing serverless frameworks helps manage spikes in AI requests.

4. Model Discovery and Lifecycle Management

- Traditional SOA might use a **service registry** to discover endpoints (e.g., UDDI or more modern API gateways).
- **AI services** add the need for **model registries**, versioning tools and performance monitoring. A new version of a model can be published and discovered similarly to any service update.

5. Impact on Traditional SOA Concepts

- **Contracts and Interfaces:** AI endpoints tend to have **input/output schemas** for inference, but can evolve over time (e.g., new features or label sets). Contract versioning remains essential.
- **Security and Governance:** Data privacy, user consent and model explainability bring new layers of governance. Sensitive data used for inference requires strict controls and audits.
- **Orchestration vs. Choreography:** AI services might be orchestrated in a **pipeline** (e.g., data cleaning → model inference → results aggregator) or use event-based triggers for asynchronous ML workflows.

6. Example Use Cases

- **Image Recognition Service:** A separate microservice wraps a CNN (Convolutional Neural Network). A front-end application calls this service via an API to identify objects in uploaded images.
- **Chatbot or NLP Endpoint:** A text analysis or large language model (LLM) is exposed as a REST or gRPC endpoint. Multiple consumer applications can tap into it for summarization, sentiment analysis, or Q&A.
- **Fraud Detection Pipeline:** Transaction data is streamed to an **AI microservice** that continuously scores for fraud risk, returning responses to a parent billing or payment service in near real-time.

Benefits and Challenges of SOA

Benefits of SOA

- **Business Agility:**
 - SOA enables organizations to respond quickly to changing market conditions and business requirements.
 - **Real Use Case Netflix:** Netflix employs SOA to continuously innovate its streaming platform. With SOA, Netflix can rapidly introduce new features, personalize recommendations and scale its infrastructure to accommodate fluctuations in viewer demand. For example, Netflix's recommendation service analyzes user preferences in real-time, leveraging microservices to deliver personalized content recommendations instantly.
- **Interoperability:**
 - SOA promotes interoperability by standardizing communication protocols and data formats.
 - **Real Use Case: Salesforce.com:** Salesforce.com leverages SOA to integrate its cloud-based CRM platform with various third-party applications and services. Through standardized APIs and web services, Salesforce enables seamless data exchange between its CRM system and other business systems, such as marketing automation tools, ERP systems and customer support platforms.

Challenges in Implementation

- **Cultural Resistance to Change:**
 - Implementing SOA often requires cultural shifts within organizations, as it may disrupt traditional development practices and organizational structures.
 - **Real Use Case: Banking Industry:** Large banks often face cultural resistance when transitioning to SOA due to the legacy nature of their systems and the hierarchical structure of their IT departments. Developers

may be accustomed to working in silos and there may be resistance from management to adopt new development methodologies. Overcoming this resistance requires strong leadership, effective communication and a focus on the benefits of SOA for delivering customer-centric solutions.

- **Complexity in Governance and Management:**

- SOA introduces complexity in governance, management and lifecycle management of services.
- **Real Use Case: Government Services:** Government agencies implementing SOA face challenges in managing service lifecycles, ensuring data security and maintaining compliance with regulations. For example, a government agency responsible for citizen services may struggle with governing access to sensitive data across multiple departments and agencies. Implementing robust governance frameworks and security policies is essential to address these challenges and ensure the integrity and confidentiality of citizen data.

Realizing the Benefits

- **Best Practices for Implementation:**

- **Real Use Case: Amazon Web Services (AWS):** AWS provides a comprehensive set of cloud services built on SOA principles. By offering a wide range of modular services, such as computing, storage and databases, AWS enables organizations to build scalable and resilient applications. Best practices include leveraging AWS services in a decoupled manner, implementing auto-scaling and fault-tolerant architectures and continuously monitoring and optimizing performance.

- **Continuous Improvement and Adaptation:**

- SOA is an iterative process that requires continuous improvement and adaptation to changing business needs and technology landscapes.
- **Real Use Case: Uber:** Uber continually evolves its platform using SOA principles to meet the demands of its global user base. By breaking down its monolithic architecture into microservices, Uber can deploy new features independently, optimize performance and scale its infrastructure

dynamically. Continuous improvement involves gathering feedback from users, monitoring system performance and iteratively enhancing services to deliver a seamless and **Microservices Architecture**

Contemporary Trends of SOA

Microservices architecture is an architectural style that structures an application as a collection of loosely coupled services, each responsible for a specific business function and independently deployable.

- **Key Characteristics:**

- **Service Decomposition:** Applications are decomposed into smaller, independently deployable services, each responsible for a specific business capability.
- **Decentralized Data Management:** Each service manages its own database or data store, enabling greater autonomy and scalability.
- **Polyglot Persistence:** Services can use different databases or data storage technologies based on specific requirements.
- **Infrastructure Automation:** Microservices rely on automation for deployment, scaling and monitoring to ensure resilience and reliability.

- **Real Use Case: Netflix**

- Netflix transitioned from a monolithic architecture to a microservices-based architecture to support its rapid growth and global expansion.
- Each microservice at Netflix handles a specific function, such as user authentication, content recommendation, billing and streaming.
- This architecture enables Netflix to scale its services independently, deploy updates faster and deliver personalized experiences to millions of users worldwide.

Cloud Computing and SOA

Cloud computing is the delivery of computing services—including servers, storage, databases, networking, software and analytics—over the internet to offer faster innovation, flexible resources and economies of scale.

- **Key Characteristics:**
 - **On-Demand Self-Service:** Users can provision and manage computing resources, such as servers and storage, without human intervention.
 - **Resource Pooling:** Cloud providers pool and dynamically allocate resources to multiple users, optimizing resource utilization and scalability.
 - **Pay-Per-Use Billing:** Users pay only for the resources they consume, enabling cost-effective and scalable solutions.
 - **Scalability and Elasticity:** Cloud services can scale up or down based on demand, ensuring performance and availability.
- **Real Use Case: Airbnb**
 - Airbnb leverages cloud computing services, such as Amazon Web Services (AWS), to power its online marketplace for lodging and tourism experiences.
 - By using cloud infrastructure, Airbnb can quickly scale its services to accommodate spikes in demand during peak booking seasons or events.
 - Additionally, cloud-based analytics and machine learning services enable Airbnb to personalize search results, recommend listings and optimize pricing for hosts.

Serverless Computing and SOA

Serverless computing is a cloud computing model where cloud providers manage the infrastructure, dynamically allocating resources as needed and users only pay for the compute resources consumed by their applications.

- **Key Characteristics:**

- **No Server Management:** Users do not need to provision, manage, or maintain servers or infrastructure, allowing for faster development and deployment.
- **Event-Driven Architecture:** Serverless applications are event-driven and respond to triggers or events, such as HTTP requests, database changes, or messages from queues.
- **Auto-Scaling:** Serverless platforms automatically scale resources based on demand, ensuring high availability and performance without user intervention.
- **Pay-Per-Use Billing:** Users are billed based on the actual resources consumed by their applications, offering cost savings and efficiency.
- **Real Use Case: Lyft**
 - Lyft utilizes serverless computing for its backend infrastructure to handle millions of ride requests and data-intensive operations in real-time.
 - By adopting a serverless architecture on AWS Lambda, Lyft can dynamically scale its backend services in response to user demand, ensuring low-latency responses and optimal performance.
 - Serverless computing enables Lyft to focus on building and improving its core ride-sharing platform without worrying about managing servers or infrastructure.

Unit-2 SOA Design Principles and Modeling

In this unit, we explore essential system design principles that form the backbone of scalable and reliable Service-Oriented Architectures.

Note: This content is inspired: Ashish Pratap Singh's comprehensive guide on system design, this will help gain practical insights into 30 key concepts.

Reference: <https://blog.algomaster.io/p/30-system-design-concepts>

YouTube: <https://www.youtube.com/watch?v=s9Qh9fWeOAK>

These concepts span core infrastructure elements like client-server models, DNS, proxies, load balancers and scaling techniques; data strategies such as database indexing, replication, sharding and caching; and modern architectural patterns including microservices, APIs, message queues and API gateways.

From a Service-Oriented Architecture (SOA) perspective, this unit delves into the essential concepts and strategies for designing services that are cohesive, granular and adaptable to change. Understanding these principles and patterns will provide insight into creating robust, scalable and maintainable service-oriented systems. These are not only vital for understanding modern distributed services and systems but also critical for succeeding in technical interviews.

Service Design Principles and Patterns

1. Service Coupling

Definition: Service Coupling refers to the degree of interdependence between any two business processes or services within a system.

Preferable State: In SOA, **weak coupling is preferred**, indicating lower dependency for increased flexibility, scalability and maintainability.

- Weak coupling allows services to evolve independently, reducing the risk of unintended consequences when modifications or updates are made.
- Services with weak coupling can adapt more easily to changes in business requirements, ensuring that adjustments in one part of the system do not propagate unexpectedly to other interconnected services.
- Example: Consider an e-commerce platform where a "Checkout Service" encapsulates functionalities such as processing payment, updating inventory and sending order confirmation emails. This service demonstrates strong cohesion by focusing on a cohesive set of operations related to completing the checkout process.

2. Service Cohesion

Definition: Service Cohesion means how closely the tasks within a service are related to perform the main task.

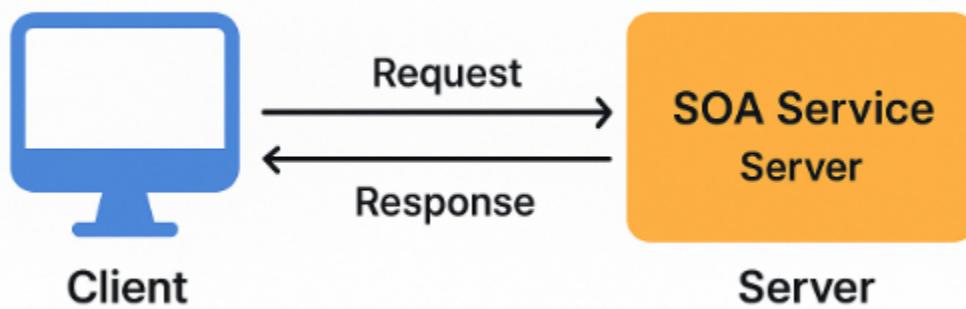
Preferable State: **Strong cohesion is preferred in SOA**, indicating that a service should encapsulate closely related and well-defined functionalities.

Explanation:

- Strong cohesion ensures that a service encapsulates a well-defined and closely related set of functionalities, enhancing clarity, maintainability and usability.
- Cohesive services promote reusability and contribute to a modular and extensible architecture.
- Example:
 - A service responsible for order processing should encapsulate functionalities such as order validation, payment processing and inventory management, exhibiting strong cohesion.

3. Client - Server Architecture

In the context of **Service-Oriented Architecture (SOA)**, the client-server architecture serves as the foundational communication model upon which services interact.



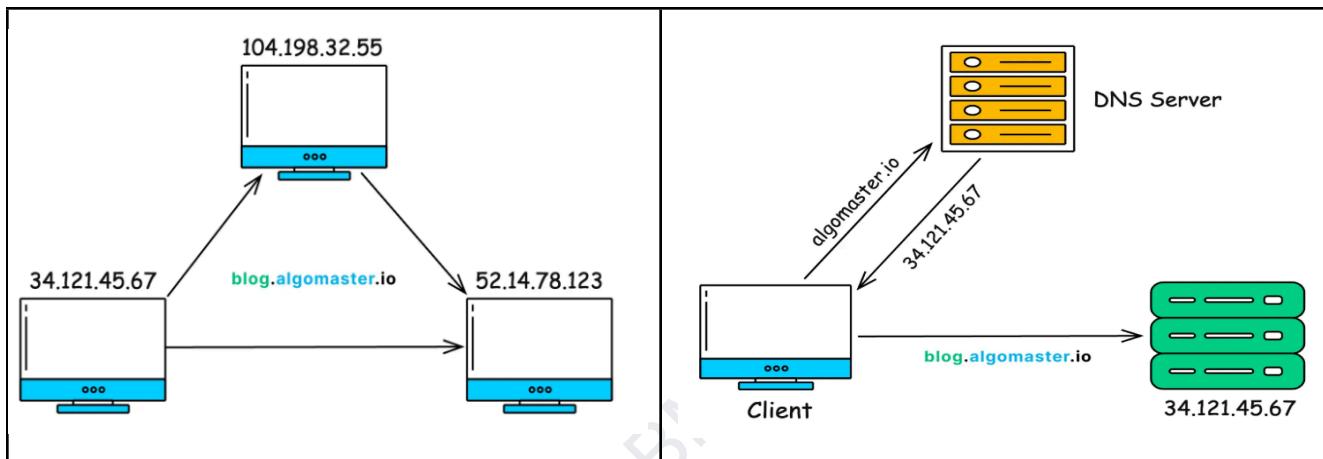
Here, the **client** can be any consumer of a service—such as a user-facing application, another service, or an external system—that initiates a request for a specific function or resource.

The **server**, in SOA, is typically represented by a **service provider**—a loosely coupled, reusable software component that exposes its capabilities through well-defined interfaces (often via protocols like HTTP, SOAP, or REST).

The client sends a request to a service to perform an operation such as retrieving data, executing a business process, or triggering an event. The **SOA service** receives this request, processes it—potentially collaborating with other services or back-end systems—and sends back a structured response.

This model allows **interoperability, scalability and reusability**, allowing different services to communicate seamlessly across platforms, technologies, or organizational boundaries.

4. IP Address and DNS

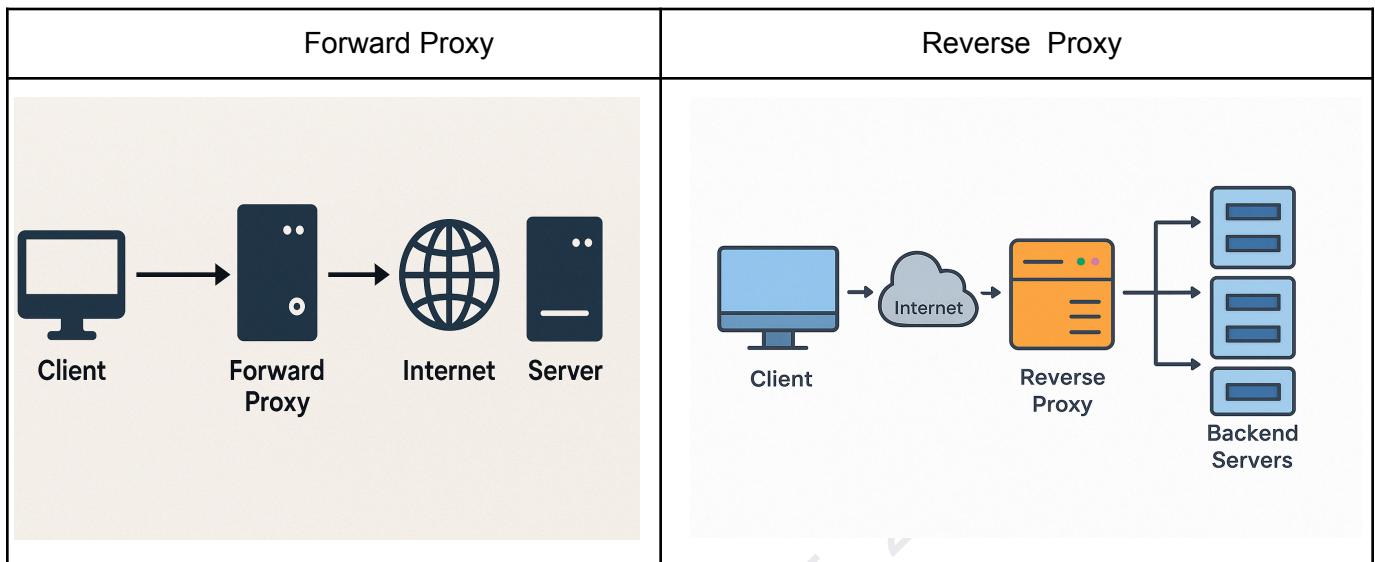


In the context of **Service-Oriented Architecture (SOA)**, the client-server architecture serves as the foundational communication model upon which services interact. Every service when deployed has a **unique IP address**. When a client wants to invoke the service, it sends service requests (via REST APIs) to the service's IP address. Example: An IPv4 address is a 32-bit number expressed in dotted notation such as : 40.1.1.1 or 192.168.22.1.

However, instead of using IP addresses, it is much easier to use invoke services using **Domain Names** that map to service IP addresses using a **Domain Name Service** or DNS protocol. Advantages of using Domain Name is that it:

1. Easy to invoke service using Domain Name (e.g. api.paytm.com)
2. Also since the IP address in the backend can change, using Domain Name makes it abstract to the client.

5. Forward Proxy and Reverse Proxy



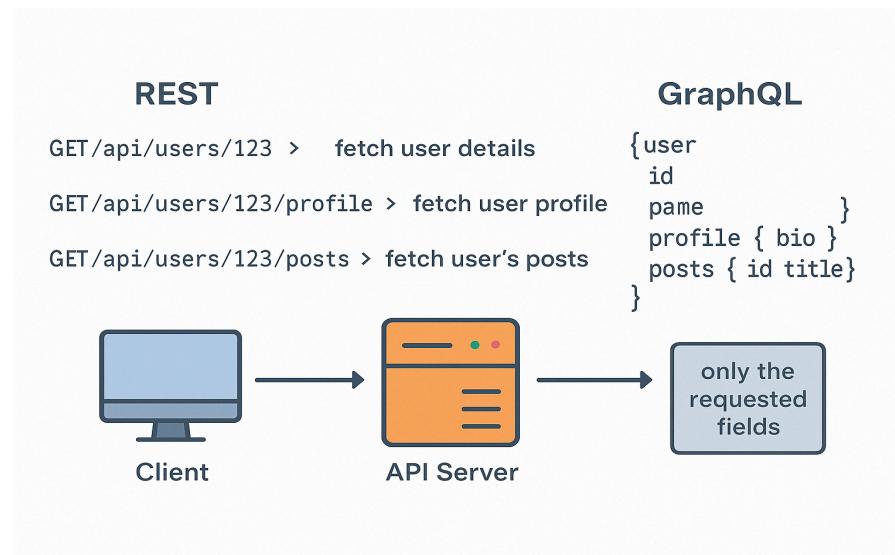
In the context of **Service-Oriented Architecture (SOA)**, proxies and reverse proxies play a vital role in enhancing **security, scalability and manageability** of service interactions.

- A **forward proxy** acts on behalf of a client, forwarding requests to external services. In SOA, this can help **secure outbound requests**, mask client identity and enforce access policies.
- A **reverse proxy**, more commonly used in SOA deployments, sits in front of backend services and **fowards requests from clients** to the appropriate backend service instance. It helps in:
 - **Securing services** by hiding their real IPs and acting as a barrier against direct exposure to security threats such as port scanning and DDoS attack.

- **Load balancing**, distributing incoming traffic across multiple service instances.
- **Routing and request filtering**, ensuring requests are handled efficiently and securely.

Overall, reverse proxies in SOA serve as intelligent gateways that provide **centralized control, traffic management** and **protection** for backend services, making the service architecture more resilient and scalable.

6. GraphQL vs REST API



In a **traditional REST** setup, retrieving specific information from the backend often means hitting multiple endpoints APIs. For instance, to get a user's basic info, their profile and recent posts, you'd typically make three separate calls:

- `GET /api/users/123` → for basic user details
- `GET /api/users/123/profile` → for profile information

- GET /api/users/123/posts → for recent posts

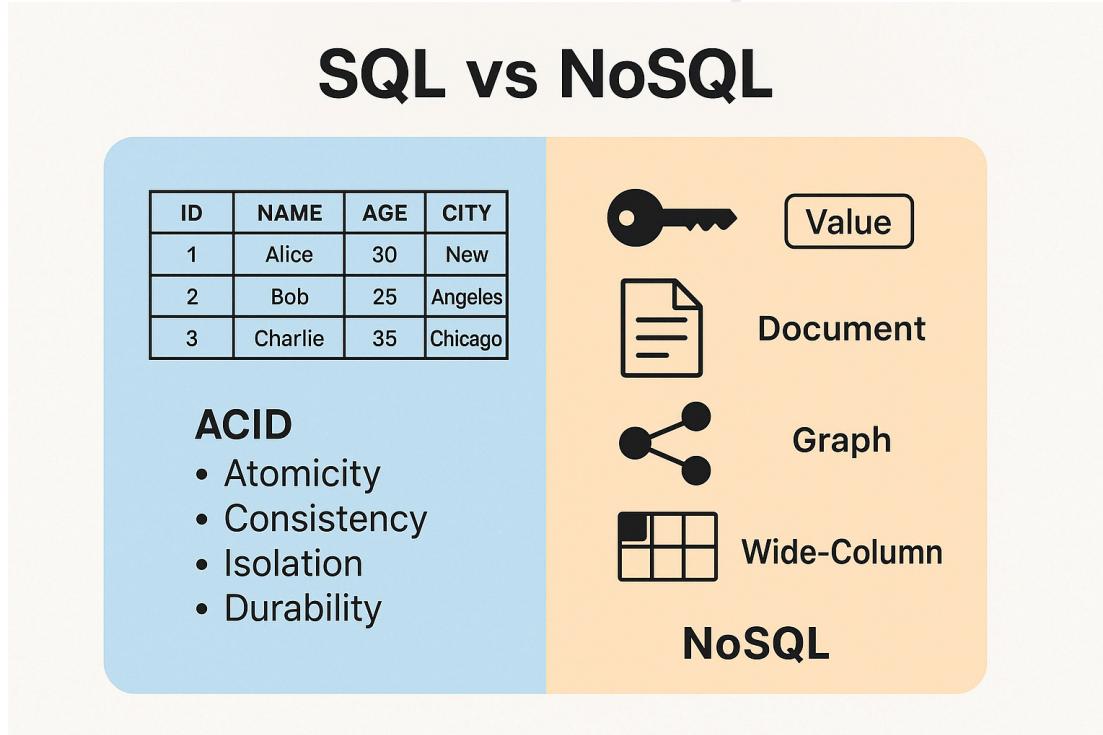
With **GraphQL**, all of this can be done in a single, structured query. Clients can specify precisely which fields they need from each related resource and the server responds with only that data—eliminating over-fetching and under-fetching.

This results in more efficient data retrieval and fewer network calls.

That said, **GraphQL also introduces some challenges**—such as more complex server-side logic and reduced caching capabilities compared to REST, which can affect performance and scalability if not managed carefully.

7. SQL vs NoSQL

In SOA-based systems, choosing the right type of data storage is crucial, as different services may have distinct data requirements.



SQL databases are commonly used for services that handle structured data with strong consistency requirements. These databases store information in relational tables and

enforce a strict schema. They adhere to **ACID properties**, which are essential in mission-critical services:

- **Atomicity**: Ensures that a service transaction is completed fully or not at all.
- **Consistency**: Guarantees that data remains valid according to defined rules after every operation.
- **Isolation**: Prevents interference between concurrent service transactions.
- **Durability**: Ensures data persists even in the event of a crash or failure.

In SOA, services like **billing, transactions, or user authentication** benefit from SQL databases (e.g., MySQL, PostgreSQL) due to their **reliability and integrity guarantees**.

On the other hand, **NoSQL databases** are better suited for services that prioritize **scalability, flexibility and speed** over rigid consistency. These databases support various models depending on the service need:

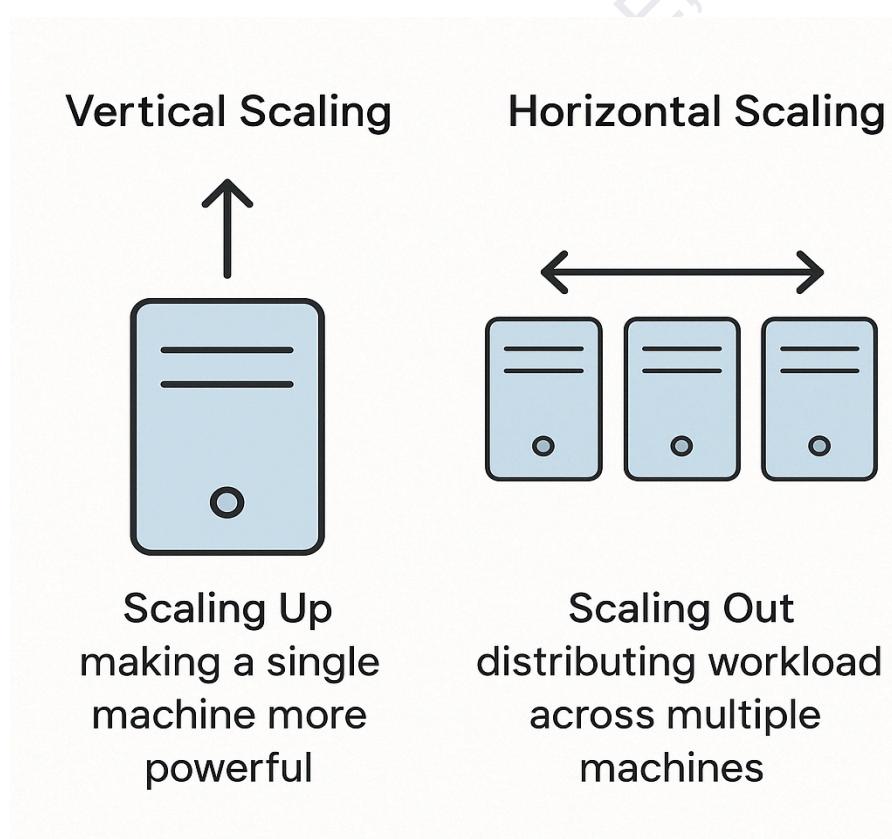
- **Key-Value Stores** (e.g., Redis) – For quick access to configuration or session data.
- **Document Stores** (e.g., MongoDB) – For storing flexible user-generated content or logs.
- **Graph Databases** (e.g., Neo4j) – For services involving relationships like recommendations or social connections.
- **Wide-Column Stores** (e.g., Cassandra) – For analytics or time-series data across distributed systems.

In modern SOA implementations, it's common to use **both SQL and NoSQL** together. For example:

- The **Order Management Service** may use SQL to ensure accurate processing and data consistency.
- The **Recommendation Service** may use NoSQL for handling large volumes of user behavior data with high read performance and flexible schema.
- This hybrid approach helps each service in the architecture **optimize performance while fulfilling its specific data consistency and structure requirements.**

8. Vertical Scaling vs Horizontal Scaling

As demand increases in an SOA-based system, the volume of service requests sent to various service components also rises. Initially, a single service instance deployed on a standard server may be sufficient. However, as more consumers start invoking services, that server can become overwhelmed, creating a bottleneck.



Vertical Scaling (Scaling Up)

A common initial response is to scale vertically—enhancing the server where the service is deployed by adding more CPU, memory, or storage resources. This makes the individual machine hosting the service more powerful.

While simple to implement, vertical scaling has significant drawbacks in a service-oriented environment:

- **Hardware Limits:** Each server has a maximum threshold. Eventually, you hit physical limits
- **Cost:** High-performance hardware becomes prohibitively expensive.
- **Single Point of Failure:** If the upgraded machine fails, all the services running on it become unavailable, affecting the entire service ecosystem

Thus, vertical scaling is a **short-term solution** and not ideal for the distributed, fault-tolerant nature of SOA.

Horizontal Scaling (Scaling Out)

To maintain scalability and resilience, SOA systems typically favor **horizontal scaling**. This involves deploying multiple instances of services across different servers or nodes. Instead of relying on a single powerful server, the load is distributed across multiple machines.

This approach aligns well with SOA principles because:

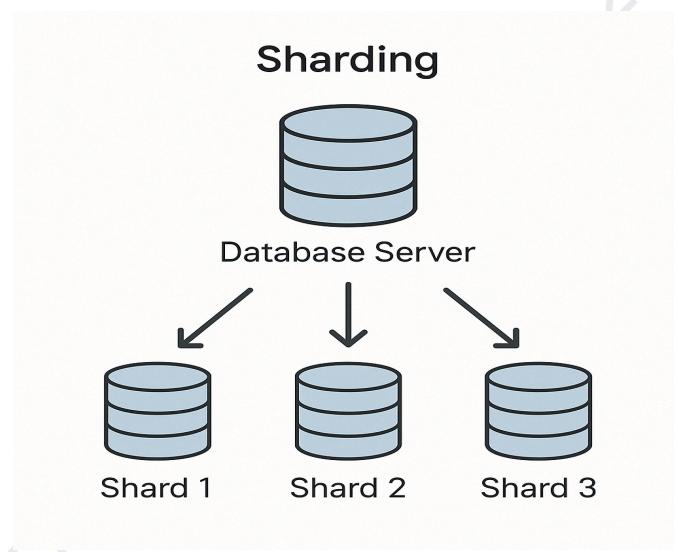
- **Improved Scalability:** More service instances across nodes can serve more client requests concurrently.
- **Fault Tolerance:** Failure of one instance or server doesn't disrupt the entire system—other service instances continue to operate.
- **Cost Efficiency:** It's often more economical to use several standard servers than a single high-end machine.

However, horizontal scaling requires **service discovery and load balancing** mechanisms. A **Load Balancer** or a **Service Registry** (like Consul, Eureka, or Kubernetes' built-in service discovery) helps route client requests to the appropriate service instance.

9. Sharding

In a large-scale SOA system with millions of users, services often rely on backing databases that grow rapidly in size—sometimes to several terabytes. As service adoption increases, a single centralized database becomes a bottleneck, affecting the performance and scalability of the entire system.

To maintain responsiveness and availability across services, SOA architects employ **sharding**—a database partitioning technique that aligns well with SOA's distributed nature.



Sharding is the practice of **splitting a large database into smaller, more manageable segments** known as **shards**. Each shard holds a portion of the data and can be hosted on a separate database server.

- Each shard contains a **subset of the overall data**, often divided by a **sharding key** (e.g., `customer_id`, `tenant_id`, or `region_id`).

- Services interacting with the database use this key to determine which shard to query.
- Since services in SOA are usually stateless and modular, this aligns well—each service can be routed to the correct shard without requiring knowledge of the entire data structure.

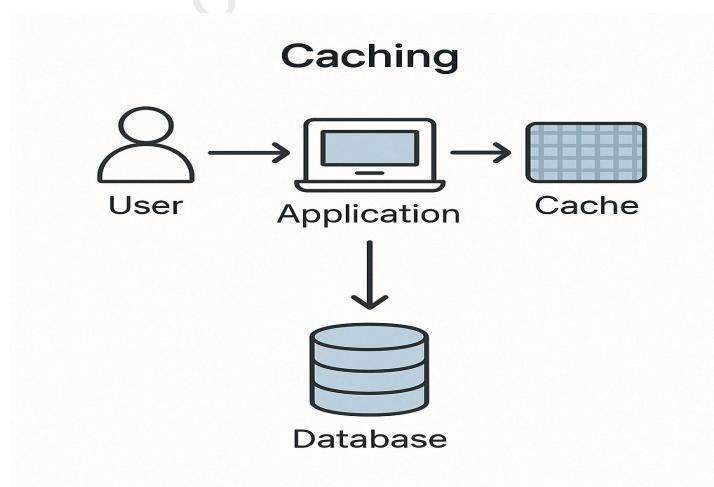
Sharding ≈ Horizontal Partitioning

- It's called **horizontal** because the data is divided **by rows**, not columns.
- For example:
 - Shard A: Users 1–1M
 - Shard B: Users 1M–2M
 - Shard C: Users 2M–3M

Each service that processes user data will interact only with the relevant shard based on the user's ID.

10. Caching

In an SOA-based system, where multiple services communicate and exchange data over the network, performance and latency are critical. One of the most effective ways to **reduce response time and minimize database load** is by implementing **caching**—storing frequently accessed data in a fast-access layer like memory.



Why Caching Matters in SOA

- Services often retrieve the same data repeatedly (e.g., product info, user preferences, lookup tables).
- Without caching, every request would trigger a database query, increasing response time and causing stress on the backend.
- Caching enables services to **respond faster, reduce repeated I/O and decrease inter-service latency**.
- **Cache Aside Pattern (Lazy Loading)**
 - One popular caching strategy in SOA is the **Cache Aside Pattern**, ideal for services that can tolerate slightly stale data.
 - Here's how it works:
 1. **Request Initiated:** A client or a downstream service requests data from a service.
 2. **Cache Check:** The service first checks if the data exists in the **cache layer** (like Redis or Memcached).
 - If **found** → Return the cached data immediately (fast response).
 - If **not found** → Proceed to the next step.
 3. **Fallback to Database:** The service queries the **underlying database**.
 4. **Cache Update:** Once data is retrieved, it's stored in the cache for future use.
 5. **Data Returned:** The service returns the data to the requester.

Handling Stale Data: Time-To-Live (TTL)

To prevent outdated or stale data from being served:

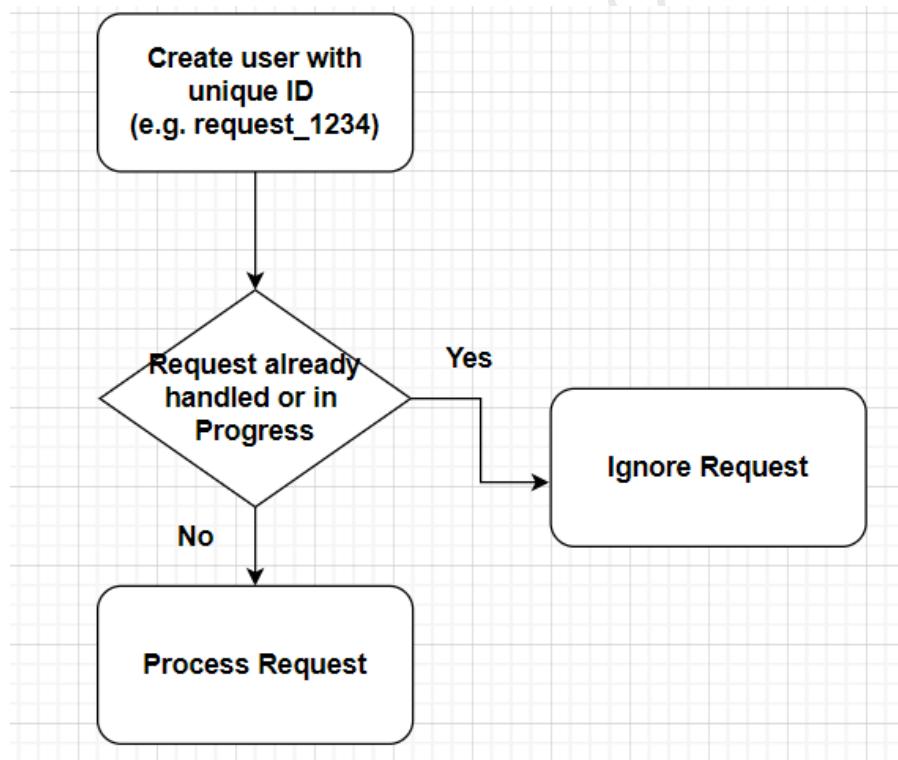
- Each cache entry is assigned a **Time-To-Live (TTL)**—a defined expiration time.
- Once the TTL expires, the cache entry becomes invalid and the service will **re-fetch updated data** from the database.

Benefits of Caching in SOA

- **Faster Service Responses** - Especially for read-heavy services and public APIs.
- **Reduced Backend Load**
Caching prevents repeated queries to the same database or downstream services.
- **Scalable Architecture**
Less stress on core services and databases improves horizontal scalability.
- **Improved Fault Tolerance**
During brief outages of backend services, cached data can still be served (depending on your freshness strategy).

11. Idempotency

In an SOA-based system, where multiple services communicate and exchange data over the network, performance and latency are critical. One of the most effective ways to **reduce response time and minimize database load** is by implementing **caching**—storing frequently accessed data in a fast-access layer like memory.



In a service-oriented architecture, where services often communicate over unreliable networks, **failures, timeouts and retries** are natural occurrences. For example, a consumer service may retry a failed request to a payment or order service. If not handled properly, such retries can result in **duplicate operations**—like charging a customer twice.

- **Idempotency** is a design principle that ensures **safety and consistency** in such distributed environments.
- Idempotency guarantees that **repeated invocations** of the same service operation with the same input will have **no additional effect** beyond the initial execution.
- This is critical in SOA, where multiple services can trigger the same action more than once due to retries or user actions (e.g., refreshing a page or re-clicking a button).

How It Works in SOA Services

1. Unique Request Identifier:

Each incoming service request is tagged with a **unique ID** (e.g., `txn-98765`), either generated by the client or the orchestration layer.

2. Duplicate Check:

Before executing the service logic, the service checks a cache or database to see if that request ID has **already been processed**.

3. Execution Path:

- If the request is **new** → The service processes it and records the result along with the request ID.
- If the request is a **duplicate** → The service **returns the stored result** or silently ignores the duplicate request.

12. Circuit Breakers

- **Definition:** The Circuit Breaker Pattern is a design pattern used to handle faults and failures in distributed systems by temporarily suspending requests to a failing service.
- In SOA, circuit breakers monitor the health of downstream services and prevent cascading failures by quickly detecting and isolating faulty components.
- **Example:** In a microservices architecture for a social media platform, a circuit breaker is implemented in the Notification Service to prevent excessive retries when sending notifications to users. If the Notification Service experiences a high error rate, the circuit breaker opens, temporarily halting requests to the service and preventing overload.

Unit-3 SOA Implementation Techniques

Web Services Standards

Web services standards define the protocols and formats used for communication between different software applications over the internet. These standards enable interoperability and integration between heterogeneous systems, allowing them to exchange data and invoke functionality seamlessly.

Simple Object Access Protocol

- SOAP is a protocol used for exchanging structured information between systems.
- It defines a standard XML format for messages, which typically include headers and bodies.
- SOAP messages are typically transported over HTTP. Other protocols like SMTP (Simple Mail Transfer Protocol) and JMS (Java Message Service) can also be used.
- SOAP provides a robust messaging framework with features such as security, reliability and transactionality.
- It follows a contract-based approach, where the structure of messages and operations is defined in a WSDL (Web Services Description Language) document.

Representational State Transfer (REST)

- REST is an architectural style for designing networked applications, emphasizing simplicity, scalability and statelessness.
- It relies on standard HTTP methods such as GET, POST, PUT, DELETE for performing CRUD (Create, Read, Update, Delete) operations on resources.
- RESTful APIs expose resources as URIs (Uniform Resource Identifiers) and use HTTP status codes for indicating the outcome of operations.
- REST APIs are lightweight, easy to understand and widely adopted for building web services, especially for public-facing APIs.

Graph QL

- GraphQL is a query language and runtime for APIs developed by Facebook.
- It allows clients to specify exactly what data they need, enabling more efficient and flexible data retrieval compared to traditional REST APIs.
- With GraphQL, clients can request multiple resources in a single query and receive only the data they ask for, reducing over-fetching and under-fetching of data.
- GraphQL APIs are introspective, meaning they expose a schema that describes the types of data available and the operations that can be performed.

Examples and Code Snippets

SOAP:

- Example: Integrating a payment gateway API into an e-commerce platform.
- **Code Snippet: SOAP**

```
<soapenv:Envelope  
xmlns:soapenv="http://schemas.xmlsoap.org/soap/envelope/"  
xmlns:web="http://www.example.com/webservice">  
    <soapenv:Header/>  
    <soapenv:Body>  
        <web:ProcessPayment>  
            <web:Amount>100.00</web:Amount>  
            <web:CardNumber>1234567890123456</web:CardNumber>  
            <!-- Additional Payment Details -->  
        </web:ProcessPayment>  
    </soapenv:Body>  
</soapenv:Envelope>
```

REST:

- Real-life Example: Retrieving weather data from a public API.
- **Code Snippet (Python using requests library)**

```
import requests  
  
url = "https://api.weather.com/data"  
params = {"city": "Bangalore", "format": "json"}  
  
response = requests.get(url, params=params)  
weather_data = response.json()  
print(weather_data)
```

GraphQL:

- Example: Fetching user profile data from a social media platform API.
- Code Snippet (GraphQL query)

```
query {  
  user(id: "123@fb.com") {  
    id  
    name  
    email  
    posts {  
      id  
      title  
      content  
    }  
  }  
}
```

Containerization and Orchestration

Containerization

Containerization is a lightweight, portable and efficient method for packaging, distributing and running applications. Containers encapsulate everything needed to run an application, including the code, runtime, libraries and dependencies, into a single unit.



The infographic illustrates the concept of containerization, showing how it connects development environments to operational stacks across various platforms and operating systems.

XENONSTACK

Containerization

Running Any App Anywhere

From Dev

- Icon: Code editor with </> symbol
- Icon: Up and down arrows
- Icon: Magnifying glass

To Ops

Any App

- Microsoft ASP.NET
- Java
- Ruby
- node
- Hadoop

More

Containerization Engine

- Icon: Gear inside a striped container box

Any OS

- Microsoft
- Linux

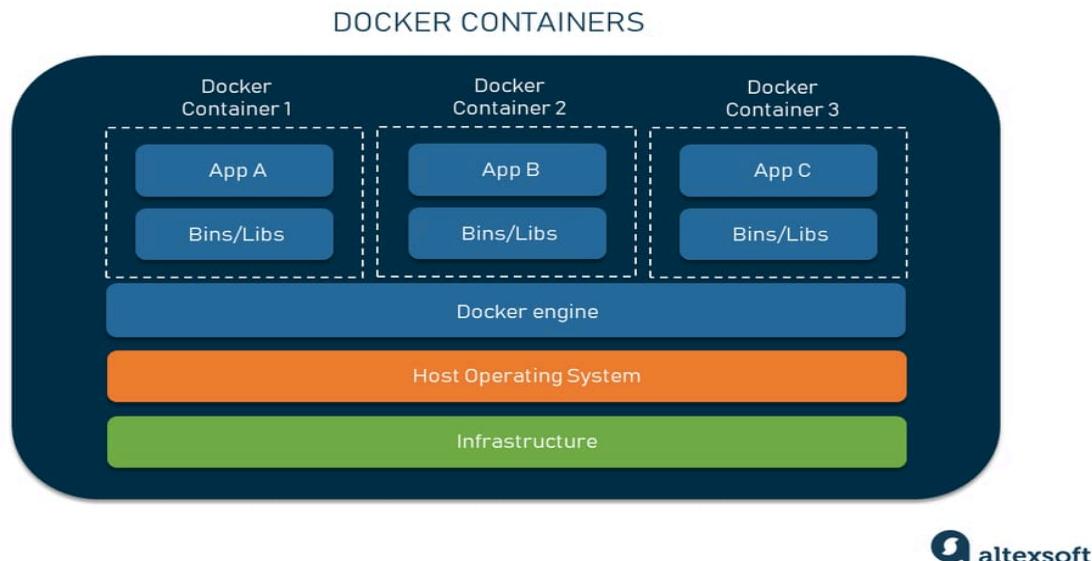
Anywhere

- Icon: Building
- Icon: Chain link
- Icon: Cloud

Source: <https://www.xenonstack.com/insights/containerization>

Docker Container

- Docker is a leading containerization platform that allows developers to build, ship and run applications in containers.
- Docker containers are isolated environments that share the host operating system's kernel, providing consistency across different environments.
- Docker uses Dockerfiles to define container configurations and Docker images to package applications and their dependencies.



 altexsoft

Source: Alexsoft

Containerization - Pros and Cons

Reference: <https://www.xenonstack.com/insights/containerization>

- Pros
 - Portability - no dependency on hardware, containers (dockers) abstracts running of application on any host
 - Lightweight - contains only application specific requirements and no unnecessary OS overhead, keep it lightweight
 - Speed - more faster and efficient in application bring up
 - Cost-effective - cost of running containers is much lower than running virtual machines
- Cons
 - Security - vulnerability of container engine and poor access control has associated risks
 - Manageability - managing large number of containers is challenging

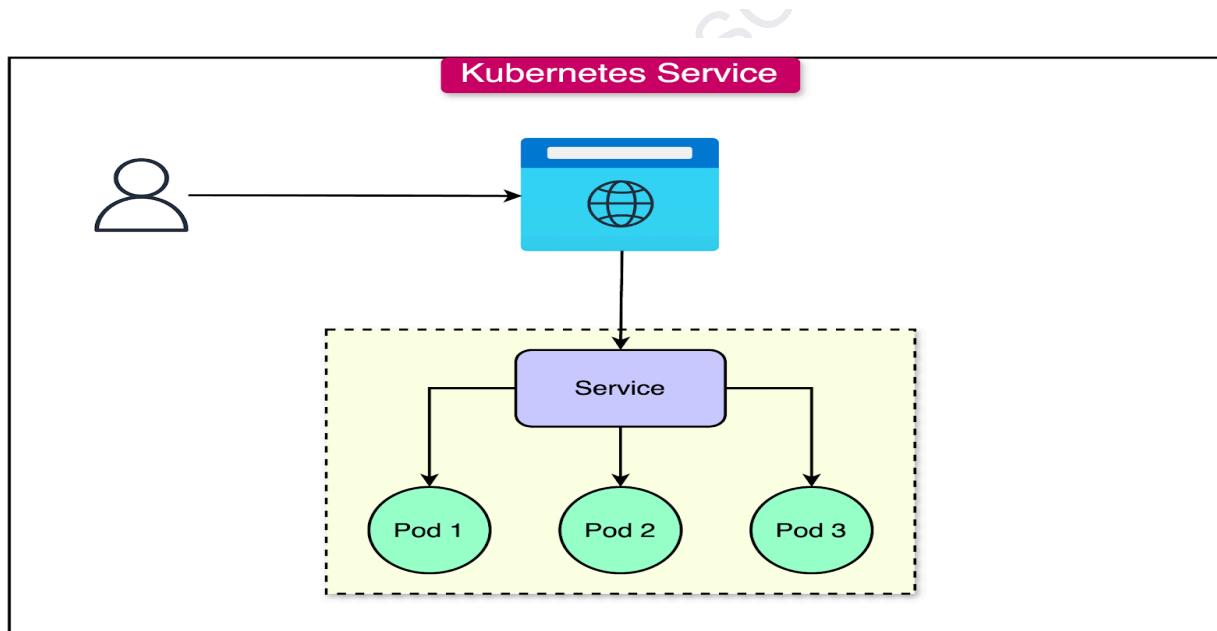
- Monitoring - needs a good monitoring system for effective maintenance and troubleshooting.

Example: A web application running in a Docker container:

```
# Dockerfile
FROM python:3.9-slim
WORKDIR /app
COPY .
RUN pip install -r requirements.txt
CMD ["python", "app.py"]
```

Kubernetes Orchestration

- Kubernetes is an open-source container **orchestration** platform for automating the **deployment, scaling and management** of containerized applications.
- Kubernetes abstracts away underlying infrastructure complexities and provides features like **automatic scaling, load balancing and self-healing**.
- Kubernetes organizes containers into logical units called **pods**, which are the smallest deployable units in Kubernetes.



Reference: <https://blog.bytebytogo.com/p/kubernetes-made-easy-a-beginners>

Example: A Kubernetes deployment manifest for a web application:

```

# deployment.yaml
# deploy 'webapp' application with 3 replicas (copies of it) access
# on port 80 where the docker image for application is:
# my-webapp:latest
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp
spec:
  replicas: 3
  selector:
    matchLabels:
      app: webapp
  template:
    metadata:
      labels:
        app: webapp
    spec:
      containers:
        - name: webapp
          image: my-webapp:latest
          ports:
            - containerPort: 80

```

Event-Driven Architecture

Event-Driven Architecture (EDA) is an architectural pattern where the production, detection, consumption and reaction to events are central to the design. EDA enables decoupled, scalable and responsive systems by promoting loose coupling between components and allowing them to communicate asynchronously through events.

Event Sourcing

- Event Sourcing is a pattern where changes to an application's state are captured as a sequence of immutable events. (**immutable: means whose value cannot be changed or will never change once set or created**)
- Instead of storing the current state of an entity, Event Sourcing stores a log of events that represent state transitions over time.
- Event Sourcing enables reconstructing the current state of an entity at any point in time by replaying the events.

Example: FinTech app

In a banking application or a FinTech application, each transaction, such as deposits, withdrawals, UPI payments, Wallet updates, is recorded as an event. The current account balance is derived by replaying these events.

Event-Driven Messaging Systems

- Event-Driven Messaging Systems facilitate communication between decoupled components by sending and receiving events.
- Event messages contain information about a specific event, such as its type, timestamp and payload data.
- Messaging systems like Apache Kafka, RabbitMQ and Amazon SNS/SQS provide reliable, scalable and fault-tolerant event delivery.

Example: Ride-sharing App

A ride-sharing application uses event-driven messaging to notify drivers of ride requests, update the status of ongoing rides and handle payment transactions.

Code Snippet: Publishing an event to a message broker (using Apache Kafka):

```
from kafka import KafkaProducer
import json

producer = KafkaProducer(bootstrap_servers='localhost:9092',
                        value_serializer=lambda v:
                        json.dumps(v).encode('utf-8'))

event = {'type': 'amount_credited', 'transaction_id': '12345',
         'amount': 100.00}
producer.send('orders', value=event)
producer.flush()
```

API Management and Governance

API Management and Governance involve the planning, design, deployment and monitoring of APIs to ensure they meet business objectives, adhere to standards and

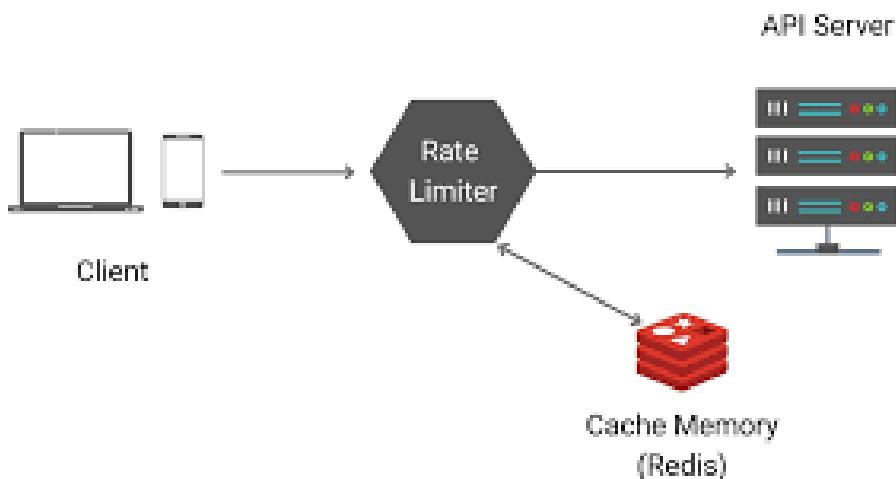
provide a positive developer experience. It encompasses various aspects such as API design, documentation, security, versioning and usage policies.

API Design Principles

- API Design Principles focus on creating APIs that are intuitive, consistent and easy to use.
- Principles include using descriptive and meaningful endpoint URLs, following RESTful design principles, using HTTP methods appropriately and providing clear and concise documentation.

Rate Limiting and Quotas

- Rate Limiting and Quotas control the number of requests an API consumer can make within a specific time frame to prevent abuse and ensure fair usage.
- Rate limits are typically enforced based on factors such as API keys, user authentication, IP addresses or subscription plans.



Source: <https://systemsdesign.cloud/SystemDesign/RateLimiter>

Example: Implementing rate limiting for a social media API to restrict users to 1000 requests per hour to prevent spamming and ensure server stability.

Code Snippet: Implementing rate limiting using Flask and Redis:

```
from flask import Flask, jsonify, request
from flask_limiter import Limiter
from flask_limiter.util import get_remote_address
from redis import Redis

app = Flask(__name__)
limiter = Limiter(
    app,
    key_func=get_remote_address,
    default_limits=["1000 per hour"]
)
redis = Redis(host='localhost', port=6379)

@app.route('/api/resource')
@limiter.limit("10 per minute")
def get_resource():
    return jsonify({'data': 'Resource data'})

if __name__ == '__main__':
    app.run(debug=True)
```

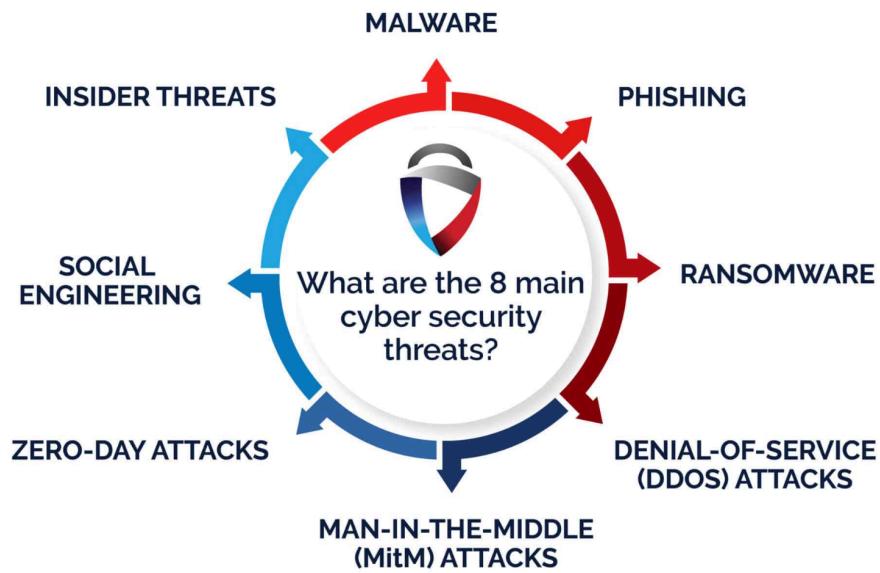
Unit-4 Security and Governance in SOA

Security Considerations in SOA

Security is a critical aspect of Service-Oriented Architecture (SOA) as it involves multiple interconnected services communicating over networks. Understanding and mitigating security risks is essential to protect sensitive data, maintain integrity and ensure compliance with regulations, such as India Data Privacy Data Protection, Europe GDPR, HIPPA and so on.

Understanding Threat Models

- Threat Models identify potential security threats and vulnerabilities that could compromise the confidentiality, integrity, or availability of services and data.
- **Common threats include** unauthorized access, data breaches, injection attacks, denial-of-service (DoS) attacks and man-in-the-middle (MitM) attacks.



Common Security Threats

Key Security Risks in SOA (Microservices) and its mitigation

1	<p>Insecure Authentication and Authorization:</p> <ul style="list-style-type: none"> • Weak authentication mechanisms like allowing weak passwords • No limit on trying wrong passwords • Providing the same privilege and permissions to all users. 	<p>Mitigation:</p> <ul style="list-style-type: none"> • Enforce Strong password policy • Limit max re-tries of incorrect username/password • Enforce Multi factor authentication (MFA), such OTP or Authenticator App • Implement Role-Based Access Control where permissions are based on roles (e.g. admin-user, standard-user, auditor-user and so on)
2	<p>Insecure Communication:</p> <ul style="list-style-type: none"> • Lack of encryption like HTTP and not HTTPS (TLS) • Use of obsolete security protocols like TLS 1.0 or MD5 hash • Improper configuration of transport layer security (TLS) can expose sensitive data to 	<p>Mitigation:</p> <ul style="list-style-type: none"> • Always use secure protocols like HTTPS (TLS), SFTP, SSH, etc • Use of TLS version 1.2 and above • Ensure TLS is properly configured to use strong ciphers and protocols (e.g. AES, SHA256, etc)

	interception to MITM (Man-in-the-Middle) attacks.	
3	<p>Denial-of-Service (DoS) Attacks:</p> <ul style="list-style-type: none"> • Overloading services with excessive requests to disrupt normal operations. • E.g. an attacker sending massive amounts of request to an API (volumetric attack) 	<p>Mitigation</p> <ul style="list-style-type: none"> • Implement rate limiting (IP address based, Username or ID based) to restrict the number of requests a client can make within a specified time frame to avoid hogging resources or over use. • Use load balancers to distribute traffic across multiple servers, preventing any single server from becoming overwhelmed. • Deploy IDPS to monitor network traffic and detect potential threat traffic.
4	<p>Data Exposure:</p> <ul style="list-style-type: none"> • Inadvertent exposure of sensitive data through misconfigured APIs or insecure storage mechanisms. 	<p>Mitigation:</p> <ul style="list-style-type: none"> • Implement strict access controls on APIs to ensure only authorized users can access sensitive data. • Encrypt sensitive data at rest using strong encryption standards • Conduct regular security audits and configuration reviews to

		identify and correct data exposure vulnerabilities
--	--	--

Example: Healthcare App

A healthcare organization's SOA faces security risks such as unauthorized access to patient records, interception of sensitive medical data during transmission between services and injection attacks targeting healthcare APIs.

key security risks specific to healthcare apps and methods to address them:

Security Risks in SOA for Healthcare Apps

1. Data Breaches and Unauthorized Access:

- **Risk:** Sensitive patient information, including personal health information (PHI), is at risk of being accessed by unauthorized users.
- **Addressing:**
 - **Strong Authentication:** Use multi-factor authentication (MFA) to ensure that only authorized personnel can access sensitive data.
 - **Access Controls:** Implement role-based access control (RBAC) to restrict access to data based on user roles and responsibilities.

2. Insecure Communication:

- **Risk:** Data transmitted between services can be intercepted, leading to the exposure of sensitive information.
- **Addressing:**
 - **Encryption:** Use TLS/SSL to encrypt data in transit. Ensure that all communications between services are encrypted.
 - **Secure Protocols:** Use secure communication protocols like HTTPS for all service interactions.

3. Injection Attacks:

- **Risk:** Improper input validation can lead to injection attacks such as SQL injection or XML External Entity (XXE) injection, compromising data integrity and confidentiality.

- **Addressing:**
 - **Input Validation:** Validate and sanitize all inputs to ensure they conform to expected formats and values.
 - **Prepared Statements:** Use prepared statements and parameterized queries to prevent SQL injection.
 - **Secure XML Parsing:** Use secure XML parsers that disable external entity resolution to prevent XXE attacks.

4. Data Exposure through APIs:

- **Risk:** Misconfigured APIs can inadvertently expose sensitive data to unauthorized users.
- **Addressing:**
 - **API Security:** Implement strict access controls and authentication mechanisms for all APIs.
 - **Data Masking:** Mask sensitive data in API responses where full exposure is not necessary.
 - **Regular Audits:** Conduct regular security audits and reviews of API configurations.

5. Denial-of-Service (DoS) Attacks:

- **Risk:** Attackers can overwhelm services with excessive requests, causing disruption of normal operations.
- **Addressing:**
 - **Rate Limiting:** Implement rate limiting to control the number of requests that can be made to a service within a specified period.
 - **Throttling:** Throttle requests to prevent any single client from overloading the system.
 - **DDoS Protection:** Use DDoS protection services to filter and manage malicious traffic.

6. Insufficient Logging and Monitoring:

- **Risk:** Lack of proper logging and monitoring can hinder the detection and response to security incidents.
- **Addressing:**

- **Comprehensive Logging:** Implement comprehensive logging of all access and activity within the system.
- **Monitoring:** Use Security Information and Event Management (SIEM) systems to monitor logs and detect anomalies in real-time.

7. Weak Service Discovery Mechanisms:

- **Risk:** Insecure service discovery can lead to the discovery and use of unauthorized or rogue services.
- **Addressing:**
 - **Secure Service Registry:** Implement secure access controls and encryption for the service registry.
 - **Authentication:** Ensure that only authenticated services can register and discover other services.

8. Compliance and Regulatory Risks:

- **Risk:** Failure to comply with healthcare regulations such as HIPAA can result in legal and financial penalties.
- **Addressing:**
 - **Compliance Audits:** Conduct regular compliance audits to ensure all services adhere to regulatory requirements.
 - **Data Protection Policies:** Implement and enforce policies for data protection, access control and data handling in accordance with healthcare regulations.

9. Integration with Legacy Systems:

- **Risk:** Legacy systems may have outdated security measures, posing risks when integrated with modern SOA services.
- **Addressing:**
 - **Secure Integration:** Use secure integration methods and protocols to interface with legacy systems.
 - **Security Patches:** Ensure that all legacy systems are updated with the latest security patches and upgrades.

10. Insufficient Data Validation and Sanitization:

- **Risk:** Poor data validation can lead to vulnerabilities such as injection attacks and data corruption.

- **Addressing:**
 - **Data Validation:** Implement rigorous data validation and sanitization processes.
 - **Frameworks:** Use secure coding frameworks that enforce data validation standards.

Security Design Patterns:

- Security Design Patterns are reusable solutions to common security problems in software architecture.
- Patterns such as these help address security concerns and reduce threats:
 - Multi-factor Authentication
 - Least Privilege Authorization Policy and Role based access control
 - Secure Communication using TLS
 - Input Validation to check for malicious input and threat injection
 - Continuous Audit Logging

SOA API Security

APIs (Application Programming Interfaces) are critical components of modern software architecture, enabling communication between different services and applications. Ensuring the security of APIs, especially in a Service-Oriented Architecture (SOA) environment, is crucial to protect sensitive data and maintain the integrity of services.

RESTful APIs are widely used in SOA due to their simplicity and scalability. Here are some best practices for securing RESTful APIs:

1. **Authentication and Authorization**
 - Use OAuth2 and OpenID for authentication.
 - OAuth2 is an industry-standard protocol for authorization. OpenID Connect builds on OAuth2 to add authentication.
2. **Data Encryption**
 - Always use HTTPS to encrypt data in transit between the client and the API server.

- This prevents eavesdropping and man-in-the-middle attacks.

3. Message-Level Encryption:

- Use message-level encryption for sensitive data within the API payload.
- Encrypt specific fields or the entire message body.

4. Input Validation

- **Sanitize Inputs:** Validate and sanitize all inputs to prevent injection attacks such as SQL injection and cross-site scripting (XSS).
- Use libraries or frameworks that provide built-in input validation.

5. Rate Limiting:

- Implement rate limiting to prevent abuse and denial of service attacks.
- Set thresholds for the number of requests a client can make in a given time period.

6. Error Handling

- Avoid exposing detailed error messages that might reveal internal server information.
- Use generic error messages and log detailed errors internally.

7. Logging and Monitoring

- **Log Requests and Responses:** Keep detailed logs of API requests and responses for auditing and troubleshooting.
- Ensure logs do not contain sensitive information.

8. Monitor API Traffic:

- Use monitoring tools to track API usage and detect unusual patterns or potential security breaches.

Unit-5 SOA Emerging Trends

Some of the emerging trends in SOA are:

1. Serverless Computing
2. AI/ML in SOS
3. Edge computing and SOA integration

Serverless Computing

Serverless computing is a cloud computing execution model where the cloud provider dynamically manages the allocation and provisioning of servers. It abstracts the server management from developers, allowing them to focus on writing code.

Serverless computing includes services like:

- Function-as-a-Service (FaaS)
- Event-driven architectures

Function-as-a-Service

Function-as-a-Service (FaaS) is a serverless computing service that allows developers to execute individual functions, pieces of business logic, in response to events without managing the underlying infrastructure.



Reference: <https://blog.back4app.com/what-are-serverless-functions-in-cloud-computing/>

Key Features

1. Event-Driven Execution:

- Functions are triggered by events such as HTTP requests, database changes, or message queue events.
- Example: AWS Lambda, Google Cloud Functions, Azure Functions.

2. Automatic Scaling:

- Functions scale automatically with the volume of incoming requests.
- No need for manual intervention to handle varying loads.

3. Pay-per-Use:

- Billing is based on the actual usage, such as the number of requests and the duration of function execution.
- Cost-effective compared to always-on server instances.

Operational Characteristics of Serverless Computing

1. No Server Management:

- Cloud provider handles server maintenance, patching and scaling.
- Developers focus on writing code.

2. Scalability:

- Automatic scaling to handle varying loads without manual intervention.
- Suitable for applications with unpredictable traffic patterns.

3. High Availability:

- Built-in redundancy and fault tolerance provided by cloud providers.
- Functions run across multiple availability zones.

4. Cost Efficiency:

- Pay only for actual usage (execution time and resource consumption).
- No charges for idle resources.

5. Quick Deployment:

- Rapid deployment of functions without lengthy provisioning times.
- Accelerates development cycles and time-to-market.

Challenges

1. Cold Starts:

- Initial latency when a function is invoked after being idle.
 - Can affect performance for latency-sensitive applications.
2. **Vendor Lock-In:**
 - Dependence on specific cloud provider services and APIs.
 - Challenges in migrating functions across different platforms.
 3. **Complexity in Debugging:**
 - Distributed nature of serverless applications complicates debugging and monitoring.

Introduction to AI and ML in SOA

Artificial Intelligence (AI) and Machine Learning (ML) are transforming Service-Oriented Architectures (SOA) by enabling services to become more intelligent, adaptive and capable of handling complex tasks. Integrating AI and ML into SOA enhances decision-making, automates processes and improves user interactions.

Artificial Intelligence (AI) and Machine Learning (ML) are increasingly being integrated into Service-Oriented Architecture (SOA) to significantly enhance service capabilities. This integration leverages the strengths of both technologies to create more intelligent, efficient and adaptive service frameworks. Here's how AI and ML are being used within SOA:

1. **Service Optimization and Automation:**
 - **Predictive Analytics:** ML algorithms analyze historical data to predict future trends and behaviors, enabling proactive service adjustments and optimizations.
 - **Automated Decision-Making:** AI systems automate routine decisions and processes, reducing manual intervention and speeding up service delivery.
2. **Enhanced Data Processing:**
 - **Natural Language Processing (NLP):** NLP capabilities enable services to understand and process human language, enhancing customer interaction through chatbots and virtual assistants.
 - **Data Integration and Management:** AI improves data integration from various sources, ensuring more accurate and real-time data availability for services.
3. **Personalization and Customization:**
 - **User Behavior Analysis:** ML models analyze user behavior and preferences, allowing services to offer personalized experiences and recommendations.

- **Adaptive Services:** AI-driven services can adapt in real-time to user needs and contexts, providing more relevant and dynamic responses.

4. Improved Security and Compliance:

- **Anomaly Detection:** AI algorithms detect unusual patterns and potential security threats, enhancing the security of the service ecosystem.
- **Compliance Monitoring:** ML models continuously monitor service activities for compliance with regulations and standards, ensuring adherence and reducing risks.

5. Operational Efficiency:

- **Resource Management:** AI optimizes the allocation and utilization of resources, improving operational efficiency and reducing costs.
- **Performance Monitoring:** ML continuously monitors service performance, identifying bottlenecks and recommending improvements.

6. Advanced Analytics and Insights:

- **Real-Time Analytics:** AI enables real-time analysis of service data, providing instant insights and enabling quick responses to emerging issues.
- **Predictive Maintenance:** ML predicts when service components are likely to fail, allowing for preemptive maintenance and minimizing downtime.

By integrating AI and ML into SOA, organizations can create smarter, more responsive and efficient service architectures that better meet the evolving needs of users and businesses. This fusion not only enhances current capabilities but also opens up new possibilities for innovation and growth in service delivery.

Intelligent Agents

Intelligent agents are autonomous entities that observe and act upon an environment to achieve specific goals. In SOA, intelligent agents can enhance services by performing tasks such as decision-making, monitoring and automation.

Characteristics

1. Autonomy:

- Operate without human intervention.
- Make decisions based on predefined rules or learned behaviors.

2. Reactivity:

- Respond to changes in the environment in real-time.
- Example: Monitoring system performance and alerting when anomalies are detected.

3. **Proactivity:**

- Take initiative to perform tasks or make recommendations.
- Example: Proactively scaling resources based on predicted load.

4. **Learning Ability:**

- Adapt and improve over time through learning from data and interactions.
- Example: Personalizing user experiences based on behavior analysis.

Example Use Case

● **Service Health Monitoring:**

- An intelligent agent monitors the health of various services in an SOA.
- Uses ML models to predict failures and automatically takes corrective actions (e.g., restarting a service, scaling resources).

Natural Language Processing (NLP)

Natural Language Processing (NLP) is a branch of AI that focuses on the interaction between computers and humans through natural language. In SOA, NLP can enhance services by enabling them to understand and generate human language.

Key Applications

1. **Chatbots and Virtual Assistants:**

- Automate customer support and interactions.
- Example: A customer service bot that answers queries and processes requests.

2. **Text Analysis:**

- Extract meaningful information from text data.
- Example: Analyzing customer feedback to identify common issues and sentiments.

3. **Machine Translation:**

- Automatically translate text from one language to another.
- Example: Translating service documentation for global users.

Example Use Case

- **Automated Customer Support:**

- Implement an NLP-powered chatbot that interacts with customers, answers common questions and escalates complex issues to human agents.

Dept. of ISE, BMSCE, 2025

Lab Exercises - Solution

Exercise 1: Overview of SOA: Implement a REST Web Service

Code Example: Develop a simple web service using a framework like Flask (Python), Spring Boot (Java), or Express (Node.js). Demonstrate how clients can consume this service to retrieve or manipulate data.

REST Web Service - Python Implementation (GET and POST Methods)

Prerequisites:

- Python Installation version 3.X
- Install Flask - use pip or pip3 based on your installation of python
- Update host firewall (if configured) to allow
- curl command

Install Flask using pip:

```
pip3 install flask
```

```
# Save the following a file called webserver.py
```

```
# Create a flask application
```

```
from flask import Flask, jsonify, request

app = Flask(__name__)

# Sample data
books = [
    {"id": 1, "title": "Book 1", "author": "Author 1"},
    {"id": 2, "title": "Book 2", "author": "Author 2"},
    {"id": 3, "title": "Book 3", "author": "Author 3"}
]

# Endpoint to get all books
@app.route('/books', methods=['GET'])
def get_books():

    # Your code here to handle GET request and return books data
```

```

        return jsonify(books)

# Endpoint to get a specific book by id
@app.route('/books/<int:id>', methods=['GET'])
def get_book(id):
    book = next((book for book in books if book['id'] == id),
None)
    if book:
        return jsonify(book)
    else:
        return jsonify({"error": "Book not found"}), 404

# Endpoint to add a new book
@app.route('/books', methods=['POST'])
def add_book():
    data = request.json
    new_book = {
        "id": len(books) + 1,
        "title": data['title'],
        "author": data['author']
    }
    books.append(new_book)
    return jsonify(new_book), 201

if __name__ == '__main__':
    app.run(debug=True)
```

```

### **Run the application using python**

```
python app.py
```

The Flask application should be running. Invoke the service using You can consume this service using HTTP client such as curl

#### **To get all books, this demonstrates GET method of Web Service**

```
curl http://localhost:5000/books
```

#### **To get a specific book by id:**

```
curl http://localhost:5000/books/1
```

#### **To add a new book:**

```
curl -X POST -H "Content-Type: application/json" -d
'{"title":"New Book","author":"New Author"}'
http://localhost:5000/books
```

#### Alternate command

```
curl -X POST -H "Content-Type: application/json" -d "{\"title\":\"New Book\", \"author\":\"New Author\"}" http://localhost:5000/books
```

## REST Web Service - Spring Boot (Java) Implementation

First, make sure you have Spring Boot installed. Prerequisites:

- Install Spring Boot:  
<https://docs.spring.io/spring-boot/docs/current/reference/html/getting-started.html#getting-started.installing>
- Install Maven: <https://maven.apache.org/>
- Install TomCat server: <https://tomcat.apache.org/download-10.cgi>
- Configure Spring Boot and Tomcat  
<https://www.baeldung.com/spring-boot-configure-tomcat>
- Add Spring Boot dependencies to Maven by creating **pom.xml**:

#### // Create pom.xml

```
<dependencies>
 <dependency>
 <groupId>org.springframework.boot</groupId>
 <artifactId>spring-boot-starter-web</artifactId>
 </dependency>
</dependencies>
```

#### // Create Java code

```
import org.springframework.boot.SpringApplication;
import
org.springframework.boot.autoconfigure.SpringBootApplication;
import org.springframework.web.bind.annotation.*;

import java.util.ArrayList;

import java.util.List;
```

```

import java.util.Optional;

@SpringBootApplication
@RestController
public class Application {

 private List<Book> books = new ArrayList<>();

 public static void main(String[] args) {
 SpringApplication.run(Application.class, args);
 }

 @GetMapping("/books")
 public List<Book> getBooks() {
 return books;
 }

 @GetMapping("/books/{id}")
 public Book getBook(@PathVariable int id) {
 Optional<Book> result = books.stream().filter(book ->
book.getId() == id).findFirst();
 return result.orElse(null);
 }

 @PostMapping("/books")
 public Book addBook(@RequestBody Book book) {
 book.setId(books.size() + 1);
 books.add(book);
 return book;
 }
}

class Book {
 private int id;
 private String title;
 private String author;

 // Getters and setters

 public int getId() {
 return id;
 }
}

```

```
public void setId(int id) {
 this.id = id;
}

public String getTitle() {
 return title;
}

public void setTitle(String title) {
 this.title = title;
}

public String getAuthor() {
 return author;
}

public void setAuthor(String author) {
 this.author = author;
}
}
```

**Run the application. Spring Boot will automatically start an embedded Tomcat server on port 8080 by default. You can now access the service using “curl” HTTP client.**

**// To get all books:**

```
curl http://localhost:8080/books
```

**// To get specific book**

```
curl http://localhost:8080/books/1
```

**// To add a new book:**

```
curl -X POST -H "Content-Type: application/json" -d
'{"title":"New Book", "author":"New Author"}
```

## **Exercise 2: Principles and Concepts of SOA**

- Code Example: Implement a basic service demonstrating loose coupling by using asynchronous messaging (e.g., RabbitMQ or Kafka). Create a publisher service that sends messages to a message broker and a consumer service that receives and processes these messages independently.

**Pub-Sub: Demonstrate a Publisher-Subscriber message exchange using RabbitMQ.**

### **Terminology**

- A **message broker** is an intermediary service that helps reliable exchange messages from one service called “producer” or “publisher” to another service called “consumer” or “subscriber”.
- **RabbitMQ** is an open-source message-broker software that originally implemented the Advanced Message Queuing Protocol (AMQP).
- One of the real time use cases of a message broker is for “**communication (comms)**” service used in banking, ticketing and e-commerce applications to post SMS or WhatsApp message to users about a transaction (e.g. credit / debit amount, order booking, shipment details, etc), In this case, a order booking service will post a message to “comms” service via message broker, this allows asynchronous and non-blocking communication between producer and consumer.

### **Prerequisites:**

- Install Erlang
  - Erlang is a programming language developed by Ericsson in 1986. Erlang is the programming language used to code [WhatsApp](#)
  - RabbitMQ is written in Erlang
  - Erlang Installation:

- Windows Installer:  
[https://github.com/erlang/otp/releases/download/OTP-26.2.3/otp\\_win64\\_26.2.3.exe](https://github.com/erlang/otp/releases/download/OTP-26.2.3/otp_win64_26.2.3.exe)
- Install RabbitMQ
  - <https://www.rabbitmq.com/docs/install-windows#installer>

## RabbitMQ tutorial - "Hello world!"

- This consists of two programs in Python; a producer (sender) that sends a single message and a consumer (receiver) that receives messages and prints them out. It's a "Hello World" of messaging.
- <https://www.rabbitmq.com/tutorials/tutorial-one-python>

### **Exercise 3: Demonstrate a Content Delivery Network (CDN)**

Design a simple Content Delivery Network (CDN) using Python with focus on distributing content efficiently to users from multiple edge servers

Tech Stack:

- Python programming language
- Flask framework (for building HTTP servers)
- Requests library (for making HTTP requests)
- Create a folder called **content** and store a short video file or an image.

#### **Step 1: Setup Edge Servers**

```
edge_server.py

from flask import Flask, send_file
import os

app = Flask(__name__)

@app.route('/content/<path:path>')
def serve_content(path):
 content_dir = 'content'
 file_path = os.path.join(content_dir, path)
 return send_file(file_path)

if __name__ == '__main__':
 app.run(host='0.0.0.0', port=5000)
```

#### **Step 2: Load Balancer**

```
load_balancer.py

from flask import Flask, request, redirect
import random

app = Flask(__name__)
```

```

edge_servers = ['http://localhost:5000', 'http://localhost:5001',
'http://localhost:5002']

@app.route('/')
def load_balancer():
 # randomly select one of the servers
 selected_server = random.choice(edge_servers)
 return redirect(selected_server)

if __name__ == '__main__':
 app.run(host='0.0.0.0', port=8000)

```

### **Step 3: Run Edge Servers and Load Balancer**

- **Open three terminal windows and run the edge servers:**

**python edge\_server.py**

- **In another terminal window, run the load balancer:**

**python load\_balancer.py**

### **Step 4: Accessing Content through CDN**

- Open a web browser and access **http://localhost:8000/content/image.jpg**
- Refresh the page multiple times to observe content served from different edge servers.

#### **Sample Results:**

- When accessing content through the load balancer, you'll notice that the requests are redirected to different edge servers randomly, demonstrating load balancing.
- Each time you refresh the page, the image file (**image.jpg**) will be served from a different edge server, showcasing content distribution.
- You can add more content to the `content` directory and access them through the CDN to observe the distribution of different content files.

This basic setup demonstrates the concept of a Content Delivery Network (CDN) using Python and Flask, focusing on load balancing and content distribution among multiple edge servers.

#### **Exercise 4: Build a AI-driven Customer Sentiment analysis service**

Design a simple AI driven Customer Sentiment analysis service using ML models and integrate it into a SOA application.

To build a simple AI-driven service using machine learning models and integrate it into a Service-Oriented Architecture (SOA), you can follow these steps and use the following tools and code snippets:

1. **Objective:** Develop a sentiment analysis service that analyzes customer reviews and provides feedback on product sentiment.
2. **Machine Learning Models:** Use a pre-trained natural language processing (NLP) model for **sentiment analysis**. For this example, we'll use the Hugging Face Transformers library with a pre-trained **BERT** mode (**Bidirectional Encoder Representations from Transformers**)
3. Refer: [https://huggingface.co/docs/transformers/en/model\\_doc/bert](https://huggingface.co/docs/transformers/en/model_doc/bert)
4. Tools:
  - Python3 for coding
  - **Hugging Face Transformers library for NLP models**
  - **Website:** <https://huggingface.co/docs/transformers/quicktour>
  - Flask for creating the web service
  - Docker for containerization
5. Implementation:
  - Install libraries: pip3 install transformers flask
  - Create a Python script for the sentiment analysis service (`sentiment_service.py`)
  - Code: save this code as **sentiment\_service.py**

```
from transformers import pipeline
from flask import Flask, request, jsonify

app = Flask(__name__)
nlp = pipeline("sentiment-analysis")

@app.route("/analyze_sentiment", methods=["POST"])
def analyze_sentiment():
 data = request.json
 text = data["text"]
 result = nlp(text)[0]
 return jsonify({"text": text, "sentiment": result["label"],
 "confidence": result["score"]})
```

```
if __name__ == "__main__":
 app.run(host="0.0.0.0", port=5000)
```

- Above script defines a Flask application with a single endpoint **/analyze\_sentiment** that accepts POST requests. When a request is received with a JSON payload containing the text to analyze, the sentiment analysis model is invoked to analyze the text and the result is returned as JSON containing the analyzed text, sentiment label and confidence score. RESTful API endpoint server is accessible via HTTP.
- Create a Dockerfile to containerize the service save it as: **Dockerfile**

```
FROM python:3.9-slim
WORKDIR /app
COPY . .
RUN pip install --no-cache-dir -r requirements.txt
CMD ["python", "sentiment_service.py"]
```

- Build and run the Docker container:  
**docker build -t sentiment-service .**  
**docker run -p 5000:5000 sentiment-service**

## 6. Integration with SOA:

- Use 'curl' to invoke sentiments API:

### Input: (Good Sentiment)

```
curl -X POST http://localhost:5000/analyze_sentiment \
-H "Content-Type: application/json" \
-d '{"text": "I like this product! It\'s awesome."}'
```

### Output:

```
{
 "text": "I like this product! It's awesome.",
 "sentiment": "POSITIVE",
 "confidence": 0.9998
}
```

### Input: (Bad Sentiment)

```
curl -X POST http://localhost:5000/analyze_sentiment \
-H "Content-Type: application/json" \
-d '{"text": "This product is bad! Don\'t buy it."}'
```

**Output:**

```
{
 "text": "This product is bad! Don't buy it.",
 "sentiment": "NEGATIVE",
 "confidence": 0.9985
}
```

Dept. of ISE, BMSCE, 2025

## Exercise 5: Contemporary Trends in SOA

Build a serverless function using a platform like AWS Lambda or Azure Functions.

Create a simple function that performs a specific task (e.g., image resizing, data processing) and expose it as a RESTful endpoint. Integrate this function into an existing SOA architecture to demonstrate its interoperability with other services

- **Objective:** The objective of this lab exercise is to create a serverless function using AWS Lambda that resizes images. The function will be exposed as a RESTful endpoint using AWS API Gateway. Additionally, you will integrate this function into an existing Service-Oriented Architecture (SOA) by making HTTP requests to the API endpoint. This exercise aims to familiarize students with serverless computing, RESTful APIs and integrating services within an SOA.
- **Tools Used**
  - AWS Lambda: A serverless computing service to run code without provisioning or managing servers.
  - AWS API Gateway: A service to create, publish, maintain, monitor and secure APIs.
  - Pillow: A Python Imaging Library (PIL) fork that adds image processing capabilities.
  - Python: The programming language used to write the Lambda function and integration script.
  - Requests Library: A simple HTTP library for Python to make API requests.
- **Prerequisites**
  - AWS Account: Access to **an AWS account** with permissions to create Lambda functions, API Gateway and IAM roles.
  - Note:
    - Try to use your college AWS account, if available, if not, create a free AWS trial account:
      - How to create a free trial AWS Account:  
<https://k21academy.com/amazon-web-services/aws-solutions-architect/create-aws-free-tier-account/>
      - !! CAUTION !!
        - ONCE THIS EXERCISE IS COMPLETE, REMEMBER TO DELETE ALL AWS RESOURCES CREATED AS PART OF THIS EXERCISE, ELSE AWS WILL CONTINUE TO "BILL" USAGE OF YOUR RESOURCE.
    - Basic Knowledge of Python: Understanding of Python programming, including handling JSON and HTTP requests.

- Basic Knowledge of AWS Services: Familiarity with AWS Lambda and API Gateway.
- Python and Pip Installed: Python 3.x and pip installed on your local machine.

- **Prerequisites**

- **Step 1: Set Up AWS Lambda Function**

1. Create the Lambda Function:

- Log in to the AWS Management Console.
- Navigate to AWS Lambda.
- Click "Create function".
- Choose "Author from scratch".
- Set Function name: `ImageResizer`.
- Set Runtime: Python 3.9 (or the latest available).
- Set Permissions: Create a new role with basic Lambda permissions.
- Click "Create function".

2. Write the Lambda Function Code:

- Install dependencies locally:

```
mkdir lambda_image_resizer
cd lambda_image_resizer
virtualenv venv
source venv/bin/activate
pip install Pillow
mkdir python
cp -r venv/lib/python3.x/site-packages/*
python/
zip -r9 function.zip python
```

3. Create your function code (`lambda_function.py`):

```
import json
import base64
from io import BytesIO
from PIL import Image

def lambda_handler(event, context):
 try:
 body = json.loads(event['body'])
 image_data = base64.b64decode(body['image'])
 target_width = int(body['width'])
 target_height = int(body['height'])
```

```

 image = Image.open(BytesIO(image_data))
 resized_image = image.resize((target_width,
target_height))

 byte_stream = BytesIO()
 resized_image.save(byte_stream, format='JPEG')
 byte_stream.seek(0)

 resized_image_base64 =
base64.b64encode(byte_stream.read()).decode('utf-8')

 response = {
 'statusCode': 200,
 'body': json.dumps({'resized_image':
resized_image_base64}),
 'headers': {'Content-Type':
'application/json'}
 }
 return response

 except Exception as e:
 return {
 'statusCode': 500,
 'body': json.dumps({'error': str(e)}),
 'headers': {'Content-Type':
'application/json'}
 }

```

- Add the function code to the ZIP file and then In the AWS Lambda Console, upload the `function.zip` file.

`zip -g function.zip lambda_function.py`

- **Step 1: Create API Gateway**

- Create a New API:

- Navigate to API Gateway.
- Click "Create API".
- Choose "REST API" and then "Build".
- **Set API name: ImageResizerAPI.**
- Click "Create API".

- Create a Resource and Method:

- Create a new resource:
  - Click "Actions" and select "Create Resource".
  - Set Resource Name: `images`.

- Set Resource Path: /**images**.
  - Click "Create Resource".
- Create a POST method:
    - With the /**images** resource selected, click "Actions" and select "Create Method".
    - Choose "POST" from the dropdown and click the checkmark.
    - In the Method Execution pane, set the Integration type to Lambda Function.
    - Select the region and enter the name of the Lambda function (**ImageResizer**).
    - Click "Save" and "OK" to give API Gateway permission to invoke your Lambda function.
3. Deploy the API:
- Click "Actions" and select "Deploy API".
  - Set Deployment stage: Create a new stage called **lab**.
  - Click "Deploy".
  - Note the Invoke URL of the deployed API
4. Step 3: Integration with SOA
1. Integrate with an Existing Service:
    - The existing service will call this API by making an HTTP POST request to the API endpoint with the image data and desired dimensions.

```
import requests
import base64
import json

def resize_image(image_path, width, height):
 with open(image_path, 'rb') as
image_file:
 image_data = image_file.read()

 image_base64 =
base64.b64encode(image_data).decode('utf-8')

 payload = {
 'image': image_base64,
 'width': width,
 'height': height
 }
```

```

 api_url =
'https://{{api-id}}.execute-api.{{region}}.amazo
naws.com/prod/images'

 response = requests.post(api_url,
data=json.dumps(payload),
headers={'Content-Type':
'application/json'})

 if response.status_code == 200:
 resized_image_base64 =
response.json()['resized_image']
 resized_image_data =
base64.b64decode(resized_image_base64)
 with open('resized_image.jpg', 'wb')
as resized_image_file:

 resized_image_file.write(resized_image_data)
 print("Image resized successfully!")
 else:
 print("Failed to resize image:",
response.json()['error'])

 resize_image('path/to/your/image.jpg', 100,
100)

```

- **Calling the Lambda Function**

- **Step-1:** Convert the Image to Base64:
  - Use a tool or a script to encode your image to base64. Here's a simple way to do it using Python:

```

import base64

def encode_image_to_base64(image_path):
 with open(image_path, 'rb') as image_file:
 image_data = image_file.read()
 return
 base64.b64encode(image_data).decode('utf-8')

encoded_image =
encode_image_to_base64('path/to/your/image.jpg')

```

- **Copy the output of this script (the base64 encoded image).**

- **Step-2: Prepare JSON payload**
  - Create a JSON payload file (`payload.json`) with the base64 encoded image and desired width and height.

```
{
 "image": "base64-encoded-image-here",
 "width": 100,
 "height": 100
}
```
- **Step-3: Run curl command**
  - **Use the following curl command to make the POST request. Replace {api-id}, {region} and base64-encoded-image-here with your actual API ID, region and base64 encoded image data.**

```
curl -X POST \

https://{api-id}.execute-api.{region}.amazon
aws.com/prod/images \
-H "Content-Type: application/json" \
-d @payload.json
```

# Reference Articles

1. **Why Amazon Retail Went to SOA Architecture**

<https://highscalability.com/why-amazon-retail-went-to-a-service-oriented-architecture/>

2. **Hugging Face:**

<https://huggingface.co/docs/transformers/quicktour>

Dept. of ISE, BMSCE, 2025