

# SOA Lab Exercises

## Exercise 1: Overview of Service-Oriented Architecture

### Objective:

Introduce students to SOA basics by creating a simple, containerized web service that clients can consume.

### Steps / Tasks

#### 1. Set Up a Basic Service

- Choose a framework (Flask/Python, Spring Boot/Java, Express/Node.js).
- Create a simple endpoint (e.g., `/products` or `/users`) that returns or manipulates data (e.g., CRUD operations in-memory or in a small database like SQLite).

#### 2. Containerization

- Create a `Dockerfile` for your service.
- Use Docker to containerize and run your service locally (e.g., `docker build`, `docker run`).

#### 3. Client Consumption

- Write a simple client script or another microservice to call your web service.
- Demonstrate basic operations (GET, POST, PUT, DELETE).

#### 4. Documentation and Testing

- Produce a small `OpenAPI/Swagger` specification to define your API.
- Use a tool like `Postman` or `curl` to test endpoints.

### Key Learning Points:

- Foundational SOA concepts (service exposure, discoverability, loose coupling).
- Introduction to containerization for easy deployment and scaling.

## Exercise 2: Principles and Concepts of SOA

### Objective:

Implement a loosely coupled service architecture using **asynchronous messaging**.

### Steps / Tasks

#### 1. Choose a Message Broker

- Use **RabbitMQ**, **Apache Kafka**, or **ActiveMQ**.
- Explain how messaging decouples the producer from the consumer.

#### 2. Publisher Service

- Create a simple service that sends messages (e.g., JSON payload) to the broker whenever an event occurs (e.g., new order placed, data update).
- Containerize it if desired (using Docker) for consistency.

#### 3. Consumer Service

- Implement a separate service that subscribes to the broker and processes messages independently (e.g., logs them, stores them in a DB, triggers a workflow).
- Ensure no direct coupling between publisher and consumer beyond the message format.

#### 4. Observability

- Introduce logging and monitoring for your publisher and consumer services (e.g., using **Elastic Stack**, **Prometheus**, or built-in broker metrics).

### Key Learning Points:

- Asynchronous communication for **loose coupling**.
- Event-driven design principles and decoupled service interactions.
- Understanding of **microservices** patterns.

### Exercise 3: Contemporary Trends in SOA

#### Objective:

Explore Serverless Computing and integrate it into an existing SOA/microservice ecosystem.

#### Steps / Tasks

1. Set Up a Simple Serverless Function
  - Pick an open source serverless platform like (**OpenFaaS, Apache OpenWhisk, Knative, Kubeless, Fission, or KEDA**).
  - Create a function that performs a specific task, e.g., image resizing, simple text processing, or a quick calculation.
2. Expose the Function as a REST Endpoint
  - Use API Gateway (AWS), Azure Function's HTTP trigger, or Cloud Functions' HTTPS endpoint to make your function externally callable.
  - Verify your function can accept inputs and return outputs.
3. Integration with Other Services
  - Invoke your serverless function from a previously created microservice or a simple client.
  - Demonstrate that the function can be part of a broader SOA. For example, upload an image via a REST service, which triggers the serverless function to resize it and then store the result in a storage service.
4. Observability and Cost Monitoring (Optional Enhancement)
  - Show how to monitor invocation counts, latency and cost metrics for your serverless function.
  - Highlight the ephemeral nature of serverless (cold starts, concurrency limits, etc.).

#### Key Learning Points:

- Basics of Function-as-a-Service (FaaS).
- Serverless integration with existing services for scalability and event-driven operations.

- Challenges like cold starts, limited runtime environment, debugging in a serverless context.

## **Exercise 4: Artificial Intelligence (AI) and Machine Learning (ML) in SOA:**

### **Objective:**

Build and integrate a **simple AI-driven service** (e.g., classification, sentiment analysis, or basic prediction) within an SOA-based architecture.

### **Steps / Tasks**

#### **1. Develop/Obtain a ML Model**

- Use a small classification model (e.g., scikit-learn or TensorFlow).
- Pre-train or load a pretrained model (e.g., for text sentiment or an image classification dataset like MNIST).

#### **2. Build a Service for ML Inference**

- Wrap the model in a REST endpoint (Flask, FastAPI, or any framework).
- Accept input data (text, image, numeric features) and return inference results.

#### **3. Containerize**

- Package the model service with **Docker** for easy deployment.
- Show how the model can be scaled independently, if needed.

#### **4. Integration and Testing**

- Integrate your AI service with a front-end client or another microservice. For instance, the client sends text or image data and the AI service returns a prediction.
- Demonstrate how updates to the model (newer version, better accuracy) can be swapped in with minimal disruption to the rest of the SOA.

#### **5. Expand with MLOps**

- Briefly mention or demonstrate how to track model versions, use a model registry (e.g., MLflow).
- Automated testing: ensure new model versions do not break the interface or degrade performance.

### Key Learning Points:

- Basic AI model serving in a **service-oriented** environment.
- Handling model versioning, data input/output formats and performance considerations.
- Implementation of MLOps

Dept. of ISE, BMSCE, 2025